

lumina-carrier — Design Document

ai-ee v1 pipeline
generated 2026-07-29 22:12:58

Contents

1	Board and Run Metadata	1
2	Overview	1
3	Requirements	1
4	Architecture	1
5	Schematic	40
6	Layout	43
7	Verification	43
8	DFM and Fabrication	43
9	Run Record (Appendix)	43
10	Artifact Index	55

1 Board and Run Metadata

Field	Value
board	lumina-carrier
workspace	boards/lumina-carrier
phase	P8
gate status	3/4 gates pass (not passing: verify)
state created	2026-07-28T13:36:22
state updated	2026-07-29T22:11:27

2 Overview

3 Requirements

requirements.md not found.

4 Architecture

Decisions of record (state.json)

Phase	What	Why
P0	Proceed P1/P2 on analyst defaults marked PROVISIONAL; H1 is the single blocking confirmation point for all 16 P0 questions	Orchestrator is a background agent with no human channel; P0 questions cannot be answered mid-run. Nothing irreversible occurs before P3/P5. SAFETY caveat: Q5 (converter isolation) and Q6 (worst-case per-rail connector current) are safety-flagged - if the human overrides either at H1, P2 architecture must be revised before P3.
P1	REJECT Skyworks Si3402-B/-C as PD controller despite the brief naming it; target TI TPS2378 (single class resistor)	Si3402-B/-C are IEEE 802.3 Type 1 only (Class 3 max; the -B revision history deleted Class 4 references). D-01 requires an at-sized stage upgradable to Type 2 by a resistor change alone - Si3402 cannot do it at any resistor value. TPS2378 states Type1/Type2 compatibility via a single RCLS (90.9R Class 3 af -> 63.4R Class 4 at, 1%). TPS2372/2373 use TWO class resistors, making the upgrade a two-part change - rejected for that reason. Consequence: AN956's '~10 W regulated' figure quoted in the briefs is a Type-1-only statement.
P1	Fine-pitch mezzanine connector families are CLOSED for the expansion connector; 2.54 mm THT is the only viable class	CAR-REQ-17 requires working voltage above 57 V. Every 0.4-1.0 mm mezzanine family JLC stocks rates 50-60 V (Panasonic AXK 60V, Hirose FX10/BM22 50V, Molex 50V, TE 50V). Only Samtec QTH/QSH clears at 175V and costs ~15 USD per mated pair against a 30 USD/carrier target. 2.54 mm CONNFLY DS1021/DS1023 rates 250V male / 600V socket = 4.4x margin, 3 A/pin single-circuit derating to 1.8 A/pin adjacent. Cost consequence: all candidates are THT and Extended - JLC hand-solder charge is unavoidable.
P1	BLOCKER for H1: provisional 15 mm board-to-board standoff is unachievable, and the board-edge RJ45 collides with the daughter board at the achievable height	Stocked 2.54 mm male (6 mm mating pin) + 8.5 mm socket mates at 11.0 mm hard-seated, ~13 mm stretched. Only a PC/104-class stackthrough reaches ~15 mm and exists only in 2x20/2x40, which would force the single-connector scheme and lose the split scheme's inherent keying. Separately, board-edge THT 8P8C magjack bodies are ~13-16 mm tall, so at an 11 mm stack the daughter collides with the jack. Options for the human: notch the daughter, low-profile jack, or set the standoff from the jack height. Recommend 11 mm standoff.

Phase	What	Why
P1	Power topology: strategy (b) separate PD interface + external ≥ 60 V buck. Lead pair TPS2378DDAR + SCT2A25STER (100 V/2 A).	Every at-capable INTEGRATED PD+converter JLC stocks is transformer-based (TPS2375x, TPS23751, MP8009A, WS3204) - strategy (a) offers no non-isolated buck at Type 2. Separate buck referenced to RTN also avoids the gate-drive/current-sense level shifting a PD controller's built-in PWM needs in a non-isolated buck (Kinetic AN162). Isolation cost is NOT the ~ 2 USD BOM delta: JLC stocks no 12 V-secondary PoE PD transformer, so isolated means an off-catalog magnetic, $\sim 2\times$ converter area, opto+TL431, and it breaks the single-JLCPCB-order assumption. Efficiency ~ 92 percent computed vs 75-77 percent for the Skyworks integrated parts - worth ~ 1.5 W of light-engine budget.
P1	CORRECTION to the earlier 'Power topology' decision record	The earlier entry was written by the orchestrator BEFORE the poe-power scout reported and without reading research/poe-power.md - the part numbers happened to match the file but the efficiency figures (92 vs 75-77 percent) and the transformer-catalog claims in that entry were not verified by the orchestrator and must not be treated as sourced. VERIFIED basis, read from research/poe-power.md section 0: recommend strategy (b) separate PD interface + external 100 V buck; TPS2378DDAR C337500 SO-8-EP 10899 stock 1.083 USD @30; SCT2A25STER C5124114 ESOP-8 3160 stock 1.614 USD @30; pair 2.70 USD @30; second source TPS2379DDAR C140293. Si3402-B AND Si3404 both disqualified (Type 1 only); Skyworks PoE+ parts effectively unstocked. Every at-capable PD IC on JLC that integrates a converter is transformer-based, so non-isolated buck at Type 2 REQUIRES the two-chip solution. Class program 90.9R (class 3, af) $\rightarrow$ 63.4R (class 4, at), one 0603 change.
P1	Q7 (module vs bare chip) closes as MODULE on cost evidence: ESP32-S3-WROOM-1-N8 C2913198. Q8 (radio functional?) is now H1-blocking because it binds the permanent outline.	Bare ESP32-S3 + external W25Q64 = ~ 4.49 USD/unit vs 4.64 USD for the module = 0.16 USD/unit, ~ 2 USD across 14 boards. The module premium does not exist at this quantity, so the low-risk choice is free. N8 chosen over N16R8 because R8 parts cap at $+65^\circ\text{C}$ ambient while N8 is $-40..+85^\circ\text{C}$, and no PSRAM keeps IO35/36/37 free. Q8 is load-bearing for MECH-02: 'radio functional' forces a permanent $\sim 6\times 18$ mm no-copper antenna keep-out at a board edge with the antenna overhanging, and demands a ≥ 0.5 A 3.3 V rail; 'radio dead' swaps to WROOM-1U-N8 C2980297 (footprint-compatible) and frees both. The outline is not reversible after P5. Antenna edge must not be shared with the shielded RJ45.

Phase	What	Why
P1	Magjack: HR911105A and all 7 other high-stock HanRun jacks are NOT PoE-capable. Only 4 PoE magjacks exist on LCSC; choice is a stock-vs-correctness tradeoff for P2.	HR911105A datasheet schematic has no power pins and LCSC parametric says Non-PoE - it is the default part everyone reaches for and it does not work here. Option A HY931147C C91754, 7693 stock, 2.148 USD @14, integrated bridge bringing V+/V- out on P9/P10, 1500 Vrms, yellow+green LEDs - stock-safe. Option B HR871150C C19724786 brings out 4 raw centre taps and is the only part publishing a rating (57VDC, 350mA per centre tap, 17.5W on 2 CTs, 35W on 4 CTs) - note 2-CT 17.5W is BELOW 802.3at's 25.5W so D-01's upgrade path needs all 4 taps, and it has only 209 pcs = 14 boards with zero spares. All three HanRun PoE jacks ARE JLC-assemblable (Assembly Type: Wave Soldering, 3.50 USD + 0.0173/joint, ~0.31 USD/board) so THT is not an assembly blocker.
P1	P5 TRAP 1 (grep-verified): rules_gen.py never reads the constraints 'voltages' key, so the 48 V 0.60 mm clearance is NOT enforced during routing	The interface-poe-48v agent grep-verified rules_gen.py: it ignores 'voltages' entirely and its power net class carries only max(fab_min, 0.2) mm clearance. Nothing makes the P7 router honour 0.60 mm; the violation only appears at P8 check_creepage, after routing is done - i.e. the expensive way. Required action at P5: add an explicit HV net class or a .kicad_dru rule carrying 0.60 mm outer / 0.10 mm inner for every 48 V net. Binding spacing source: IPC-2221B B2/B1 51-100 V band (48 V and 57 V give the same answer), corroborated by TI's 0.635 mm VSS-to-VDD PD layout note. The coated column B4 (0.13 mm) is refused because check_creepage implements only B2/B1, so relying on soldermask fails P8.
P1	P5 TRAP 2: the magjack plane void must be built from plane 'region' rectangles because planes_gen has no keepout key, or P8 check_return_path fails unwaivably	check_return_path exits 2 on an absent high-speed net and ERRORS (gate.py offers no waiver) if the MDI pair crosses the deliberate plane void under the magjack. planes_gen has no keepout key, so the void cannot be declared directly - it must be produced by sizing the plane region rectangles at P5. Also recorded: 4-layer JLC04161H-3313 is REQUIRED, being the only stackup with a diff_100 profile (0.260/0.210); 2-layer computes to 1.081 mm per leg at a clamped 0.30 mm gap. Gate-5 answer: W5500 SPI binding maximum is 33.3 MHz guaranteed, NOT the 80 MHz 'theoretical design speed' - recommend 20 MHz on SPI2 IO_MUX pins. Do NOT fit 75 ohm Bob Smith resistors (TI SNLA079D 2.3: does not apply for PoE).
P2	Stackup: 4-layer JLC04161H-3313, 1 oz	Forced by 100 ohm MDI - it is the only stackup carrying a diff_100 profile (0.260/0.210); a 2-layer board computes to 1.081 mm per leg, which is unroutable. Reinforced by plane discipline (continuous In1 GND under the MDI corridor) rather than by thermals.

Phase	What	Why
P2	MECH-02 outline proposal: 100 x 80 mm, 3 mm corner radius, 4x M3 at 5 mm inset plus a 5th M3 at (46,74); 11.0 mm stack height	Derived from 4410 mm ² of measured block area at 55-60 percent realistic utilisation for a 4-layer board carrying a 0.60 mm creepage envelope, a no-vias MDI corridor, an antenna keepout and five mounting-hole exclusions. 80x60 would need 92 percent utilisation and is impossible. The 5th hole is CAR-REQ-15 mechanical support near the connector. 11.0 mm is what stocked 2.54 mm parts actually mate at; the provisional 15 mm is unachievable.
P2	ICD-01 FROZEN (draft, pending H1): split connector, 38 positions - POWER J3 2x7 + SIGNAL J4 2x12, CONNFLY DS1021/DS1023 2.54 mm	Split beats a single 2x20 on mis-mating consequence: a one-position offset on a single 2x20 puts 48 V onto an ESP32 GPIO, whereas on a 2x7 power block it lands on 12 V or GND. The asymmetric 2x7 + 2x12 pair also self-keys, which matters because MECH-01's symmetric 4x M3 pattern would otherwise allow a daughter to be bolted 180 degrees rotated. Derating 1.80 A/pin (3 A single-circuit x 60 percent adjacent). GND is the binding rail at 7 pins in the power block - the brief's '>= 4 GND' is below what CAR-REQ-13 requires.
P3	H1 approved; Q8 Wi-Fi FUNCTIONAL overrides carrier brief section 6	Owner deliberately overrode 'wireless out of scope, radio is a debugging fallback'. Carrier now carries: WROOM-1 (PCB antenna, NOT the 1U u.FL variant) at a board edge with the antenna keepout honoured on EVERY layer - no copper, no pour, no plane; RF review in scope at layout sign-off; the permanent P5 outline must allocate antenna clearance up front. This promotes the stackup.md s7.1 P5 re-basing step from mandatory-by-quality to mandatory-by-requirement: skipping it puts the antenna over solid copper and the radio is now a functional requirement. Owner also flagged an out-of-scope firmware consequence: the 00 s3 control contract (UDP 5568, 60 fps, <100 us) is specified over Ethernet; a functional Wi-Fi path either honours it with worse jitter and no PoE, or is scoped to commissioning only.
P3	RULING: no carrier-side ID EEPROM. Removed M24C02-RMN6TP (C83836) from parts.json; 47 -> 46 lines.	The frozen ICD-01 and blocks.md both put the ID EEPROM on the DAUGHTER, riding the shared I2C bus - which is exactly why ID is one pin (ID_ADC divider) and not two. My own P3 amendment 4 asked for a carrier-side EEPROM in error; the architecture and the frozen ICD are authoritative and the part-sourcer was right to flag the conflict rather than silently fit it. CAR-REQ-07 is satisfied without it (divider on ID_ADC for cheap daughters, I2C EEPROM on the daughter for the par/strobe). Carrier-side storage needs are already covered by the WROOM-1-N8's 8 MB flash and the ESP32 eFuse MAC. Deleting one line changes nothing else: no net, no connector pin, no sheet.

Phase	What	Why
P3	Open item B CLOSED from the datasheet: HY931147C supports Mode A + Mode B + either polarity	Part-sourcer read the HY931147C 2-page datasheet schematic (LCSC wmsc mirror; the normal CDN URL returns HTML). The jack contains TWO full diode bridges - one fed from both transformer centre taps (Mode A, data pairs), one from the tied spare pairs (J4+J5, J7+J8, Mode B) - both rectifying onto a common V+ (P9) / V- (P10). Mode A+B is mandatory for an 802.3-compliant PD and is present, so this is not a blocking escalation. P4 also needs: P7/P8 are absent from the schematic (no-connect on this SKU, unlike HR861153C), and LED anodes are pins 11 (yellow) / 13 (green) with cathodes on 12 / 14, so the W5500's active-low LED outputs drive the CATHODES.
P3	Open item A STANDS as a documented limitation: the selected magjack is 802.3af-qualified only. D-01's at upgrade carries TWO non-board dependencies.	HY931147C's entire published electrical spec is isolation 1500 Vrms, OCL 350 uH @ 8 mA DC bias, insertion/return loss, CMR, crosstalk. There is NO current, power or thermal rating for the internal bridge or the centre taps anywhere in the datasheet, and its only PoE claim is 'Meets or Exceeds IEEE802.3af standards including 350uH Min OCL with 8mADC'. So it is qualified for build 1 (af, Class 3) and NOT qualified for the at upgrade. Not resolvable by sourcing: no LCSC PoE magjack publishes an at (600 mA/pair) tap rating; closing it needs HanRun's at data or a Bel/Pulse PoE+ ICM off-LCSC, which breaks the single-JLPCB-order assumption. HR871150C stays rejected (its own published 17.5 W two-centre-tap figure is below 25.5 W, and 209 pcs is 14 boards with zero spares). Net effect: D-01's 'resistor change, no respin' remains true at BOARD level, but the at upgrade also requires magjack qualification AND enclosure ventilation. Escalate at H2.
P3	CBULK substituted: 4 x 10uF/100V X7R 1210 replaces the architecture's 2 x 22uF/100V	No 22 uF / 100 V MLCC exists on LCSC in any package (-package 1210 and 1812 both return zero; the only hits are aluminium electrolytics, which this board bans on the PD rail). 40 uF nominal derates to ~20-24 uF under DC bias at 48 V - still 4x the >=5 uF AC-MPS floor and far below the ~180 uF 802.3 port ceiling, so both compliance bounds hold. It is the most expensive passive line on the board at 2.04 USD/assembly. Also recorded: L20 = FXL1360-680-M was selected by a dissipation budget, not a catalogue preference - only its 140 mOhm DCR closes blocks.md's hard 1.25 W ceiling for the 48->12 block (U20 0.43 + D20 0.54 + L20 0.245 = 1.22 W); every other stocked 68 uH part is 190-260 mOhm. Cost: a 13.5 x 12.6 mm body, not the ~7x7 mm blocks.md assumed - P6 placement must budget the real footprint.

Phase	What	Why
P3	BOARD-KILLER CAUGHT AT P3: SCT2A25 feedback reference is 1.2V, not the 0.8V the architecture assumed. The sourced +12V divider (140k/10k) would produce 18V.	Datasheet states 1.2 V +/-1% (1.188/1.2/1.212) in four independent places: features list, pin-function table, EC table VREF, block diagram. Architecture blocks.md and the part-sourcer both assumed 0.8 V. 140k/10k on a 1.2 V reference = $1.2 \times 15 = 18.0$ V on the +12V rail, which would exceed the rating of every downstream part including the 22uF/25V bulk and the TPS563201 input. Datasheet Table 2 gives the vendor's own values for 12 V: $R2/R1 = 271k/30k$ plus a 150 pF feedforward cap across the top leg, which the datasheet calls NECESSARY for loop stability. NONE of 271k, 30k or 150pF exist in parts.json - this is a SOURCING gap, not a value substitution. Must be closed before P4 wires the divider. The part-sourcer had correctly flagged R60/R63 as placeholders pending extraction, so the process caught it; the extraction is what closed it.
P3	P4 design constraints from the SCT2A25 extraction: EN abs max 6V vs a 48V-referenced driver; asynchronous topology; TM pin restriction	1) EN absolute maximum is 6 V but the architecture drives it from U1 (TPS2378) CDB. If CDB swings above 6 V it destroys the buck - P4 must fit a divider or clamp, and that divider doubles as the programmable UVLO ($VIN_{rise} = 1.25 \times (R3+R4)/R4$, $VIN_{hys} = 2.1uA \times R3$). EN floats HIGH, so an absent driver turns the converter ON. 2) The part is an ASYNCHRONOUS buck - an external catch diode SW->GND is REQUIRED (SS510 is already in the BOM, correct). 3) TM pin 6 is factory-test only and must be tied to EN, tied to GND, or left floating - never wired to an arbitrary net. 4) Fsw is fixed 300 kHz, no RT pin. 5) SW abs max carries a -1 V negative floor, which with an async topology is the binding constraint on catch-diode forward drop plus ringing. 6) No recommended land pattern is published (outline only), so the footprint must come from an IPC SOIC-8+EP generator and fp_verify will have no grounded pad sizes for this part.
P3	W5500 grounded: single 3.3V rail (no 1.8V core, no RSET_BG). Two premises in the orchestrator's own extraction prompt were wrong and the agent corrected them.	I wrote '1.8 V core supply' and 'RSET_BG' into the assignment from memory of the part family. Neither exists. Truth from WIZnet W5500 DS v1.1.0: ONE 3.3 V supply feeding two on-die domains (AVDD analog pins 4/8/11/15/17/21, VDD digital pin 28; 2.97/3.3/3.63 V); the core is generated on-chip and 1V2O pin 22 is a 1.2 V internal-regulator OUTPUT needing a 10 nF cap; the bandgap-bias resistor is EXRES1 pin 10, 12.4 kohm 1% to AGND, and VBG pin 18 must be left FLOATING. The agent refusing to fill from memory is the behaviour that makes this pipeline safe - a wrong pin here is a dead board. Note EXRES1's 12.4k is the same value already in the BOM for the PD detection split, so it is a quantity change, not a new line.

Phase	What	Why
P3	P4 W5500 wiring facts: RSVD pin asymmetry, mandated caps, SPI ceiling, no exposed pad	RSVD ASYMMETRY - pin 23 must be tied to GND while RSVD pins 38-42 are NC (per the v1.0.2 correction). Easy to get wrong and a P4 trap. Pin 7 DNC is typed AI/O in Table 2 but described 'Do Not Connect'. MANDATED CAPS - only two are actually specified by the datasheet: TOCAP pin 20 needs 4.7 uF on a SHORT trace, 1V2O pin 22 needs 10 nF. The datasheet contains NO decoupling requirement and NO layout chapter at all (the words decoupling/bypass/layout appear nowhere in 67 pages), so the 100nF-per-power-pin entries are labelled NOT SPECIFIED BY THIS DATASHEET and must not be treated as grounded. SPI - FSCK is quoted '80/33.3 MHz'; 33.3 MHz is the binding guaranteed number and 80 MHz is explicitly 'theoretical design speed'. The planned 20 MHz is legal with 1.7x margin, closing schematic sign-off gate 5. Mode 0/3, MSB first. PACKAGE - LQFP-48 7x7 body, 0.50 BSC, JEDEC MS-026 BBC, NO exposed pad (the centre circles are mould/ejector holes). No recommended land pattern published, so use IPC-7351 LQFP-48. MDI pins TXN=1 TXP=2 RXN=5 RXP=6; magnetics 1:1, 350 uH; NO auto-MDIX. No MDI termination values anywhere - they must come from the magjack datasheet.
P3	ESP32-S3-WROOM-1-N8 grounded. GPIO budget CONFIRMED wider than feared: IO35/36/37 ARE available on -N8; GPIO22-34 do not exist on any pad.	Datasheet footnote b withholds IO35/36/37 only on R8/R16V Octal-PSRAM SKUs - the no-PSRAM N8 exposes them, which is one reason the N8 was the right pick. GPIO22-GPIO34 appear nowhere in Table 3-1 (not bonded out). The internal SPI-flash bus (SPI-CLK/CS0/D/Q/WP/HD) is on-module and never exposed, and the datasheet never prints its GPIO numbers - so the common '26-32 are reserved' claim was NOT asserted from memory. USB pins GPIO19 (D-) / GPIO20 (D+) are NOT dedicated and are freely usable as ordinary GPIO when USB is unused, which materially relieves open item E's 28/28 pin budget. Supply: 3V3 pin 2 sole supply, external supply must deliver min 0.5 A (peaks 355 mA Wi-Fi TX), 22 uF + 0.1 uF decoupling - matches the power tree. EN needs 10k/1uF RC, never floating, tSTBL >= 50 us. Ambient -40..+85 C for -N8, measured OUTSIDE the module - relevant to the architecture's 56/69 C internal-air figures.

Phase	What	Why
P3	P4 BRICK RISK: GPIO45 must not be pulled high at power-up on the N8, and GPIO3 has no internal pull	Section 8: on this no-PSRAM SKU a high GPIO45 at power-up forces VDD_SPI to 1.8 V and the 3.3 V on-module flash will not boot - the cheapest way to brick the design, and it is a passive-component error, not a firmware one. Strapping defaults per Table 4-1: GPIO0 weak pull-up = 1 (SPI Boot), GPIO46 weak pull-down = 0, GPIO45 weak pull-down = 0, and GPIO3 is FLOATING with no internal pull at all - if strapped it must be driven by a non-Hi-Z circuit. Strapping hold time tH is min 3 ms after EN goes high, which constrains the EN RC. P4 must verify no PWM/SPI/I2C assignment lands on 0/3/45/46 with a pull-up, and the schematic reviewer must check this explicitly.
P3	ANTENNA CLEARANCE CONFLICT: the architecture's 10x22 mm keepout is smaller than Espressif's base-board guidance now that Wi-Fi is functional	Datasheet gives only the MODULE-side 18x6 mm Antenna Area. The base-board rule lives in the separately-fetched ESP32-S3 Hardware Design Guidelines (Fig 26, tagged [HDG] in the extraction): the antenna should overhang the board edge ENTIRELY; if it does not, notch the board 6 mm deep with ≥ 15 mm LATERAL clearance on the feed-point side, ≤ 2 mm board on the far side, ≤ 1 mm overlap. Housing clearance ≥ 15 mm in all directions. The P2 architecture declared a 10 x 22 mm keepout at the right edge, which does not obviously satisfy the ≥ 15 mm lateral figure. H1 closed Q8 as 'Wi-Fi functional, a supported control path', so this is now a functional requirement rather than a nicety, and the 100x80 outline is permanent. Must be reconciled at P5/P6 placement and raised at H2. Also note: 'no copper on ALL layers' is NOT stated in those words for the PCB antenna in either PDF - the HDG's stronger rule is to remove BOARD MATERIAL in the clearance area; treat 'no board material in the clearance zone' as binding and do not cite a phrase that does not exist.

Phase	What	Why
P3	RESOLVED: HV clearance raised 0.60 -> 0.635 mm. TI's TPS2378 layout guidance exceeds the IPC-2221B number the board was using.	TPS2378 datasheet sec 7.3.9: 'TI recommends maintaining a clearance of 0.025 in between VSS and high-voltage signals such as VDD' = 0.635 mm. The board's 0.60 mm came from IPC-2221B B2/B1 (51-100 V band) and is BELOW TI's part-specific recommendation around U1. Rather than waive a vendor recommendation on the most safety-critical IC, adopt 0.635 mm as the board-wide HV clearance - the delta is 0.035 mm and costs essentially nothing in routing area. This supersedes the 0.60 mm figure everywhere it appears (constraints.json notes, stackup.md, connector-icd.md). IMPLEMENTATION: P5 TRAP-1 already requires a hand-written .kicad.dru rule because rules_gen ignores the voltages key - that rule must now carry 0.635 mm, not 0.60 mm. Note also drawing note 9: 'Size of metal pad may vary due to creepage requirement.'
P3	P4 FIX REQUIRED: TPS2378 pin 8 (APD) must be tied to RTN, not left unconnected as parts.json states. Exposed pad is VSS, never RTN.	parts.json's role note says 'Leave pin 8 unconnected'. The datasheet sec 7.3.1 says explicitly: 'If not used, connect APD to RTN.' APD sinks only ~1.73 uA, so a floating APD can drift; above 1.5 V over RTN it turns the pass MOSFET OFF, disables CLS and forces T2P - a dead board from an unconnected pin. Fit APD to RTN, optionally through a 0 ohm/DNP link to preserve the TPS2379 second-source option (its pin 8 is GATE, not APD). SEPARATELY, the exposed PowerPAD is internally connected to VSS (sec 7.3.9) - VSS is the rectified-PoE NEGATIVE rail, while RTN (pin 5) is the switched load return. Wiring the pad to RTN or to converter ground destroys the part. CDB and T2P are open-drain referenced to RTN, not VSS - another easy P4 error. MPS is NOT automatic: valid MPS needs ≥ 10 mA DC AND AC impedance below $26.3k \parallel 0.05$ uF; the AC half is met by CBULK ≥ 5 uF, the DC half must be held by the application load, so firmware must never let the board idle below 10 mA.

Phase	What	Why
P3	D-01 class lever CONFIRMED against the datasheet's own Table 1; KiCad 10 stock footprint matches TI's current EP drawing	TPS2378 Table 1 / sec 7.3.3: Class 0 = 1270R, 1 = 243R, 2 = 137R, 3 = 90.9R, 4 = 63.4R. Board fits 90.9R (Class 3, 6.49-12.95 W) and the documented upgrade is 63.4R (Class 4, 12.95-25.5 W). Sec 7.4.5 is explicit that 63.4R is the ONLY valid type-2 signature and 90.9R makes it a compatible type-1 device - so D-01's resistor-only lever is vendor-documented, not inferred. FOOTPRINT: the LCSC-hosted PDF's mechanical pages are image-only and carry the SUPERSEDED drawing 4208951-6/D; TI's own SLVSB99C carries the current 4214849/B (09/2025) and the EP grew from 3.10x2.40 to 3.4x2.71 between revisions. The agent read the KiCad .ki-cad_mod to confirm that Package_SO:SOIC-8-1EP_3.9x4.9mm_P1.27mm_EP2.95x4.9mm_Mask2.71x3.4mm matches TI's current EP exactly - use the stock KiCad footprint, and make sure the librarian does not pull a stale EasyEDA one instead.
P3	CAR-REQ-08 BLOCKER RESOLVED: TPS16630 has NO EN pin - only active-low SHDN with an INTERNAL PULL-UP, so it defaults ON. External pull-down is MANDATORY.	I asked the extractor to say loudly if the part defaults ON with a floating enable. It does. V(SHDN) open-circuit is 2.48 V min / 2.7 V typ, and the enable threshold V(SHUTR) is 2.0 V rising - so a floating SHDN sits ABOVE the enable threshold and +48V_SW goes live with the MCU absent, held in reset, or tri-stated mid firmware update. That is a direct violation of CAR-REQ-08 (de-asserted by default, asserted by firmware only after successful boot). FIX: external pull-down on SHDN, sec 8.3.11 requires ≥ 10 uA sink capability while holding < 0.8 V (I(SHDN) leakage at 0 V is -10 uA max). The 10k 0603 already in the BOM gives ~ 0.1 V against the internal pull-up with 80 uA of sink capability - ample. Polarity works out correctly: pull-down = default OFF, MCU drives SHDN HIGH to enable. SHDN abs max is 5.5 V so a 3.3 V GPIO is fine and the rail is NOT tolerant - never wire it to 48 V. MODE pin must be left OPEN for latch-off (shorted to GND = 648 ms auto-retry).

Phase	What	Why
P3	MARGINAL + BOM ERROR: eFuse dV/dt cap is 1nF against a 10nF datasheet MINIMUM, and the 2800uF daughter load charges in ~134ms against a 162ms latch-off timer	TWO issues from the same interaction. (1) BOM ERROR: parts.json fits 1 nF as the U22 dV/dT cap, deliberately 'sized FAST'. Datasheet states C(dVdT) MINIMUM 10 nF - 1 nF is out of spec and must be corrected at P4. (2) TIMING: with the daughter's ~2800 uF on +48V_SW, dV/dt ramping cannot hold inrush under the 1.0 A ILIM (at C_dVdT=10 nF, $t_{dVdT} = 20.8e3 \times 48 \times 10e-9 = 9.98$ ms, implying $I_{inrush} = C \times V / t = 13.5$ A). The eFuse therefore rides in CURRENT LIMIT for the whole charge: $t = C \times V / I = 2800e-6 \times 48 / 1.0 = 134$ ms, against a 162 ms latch-off timer with MODE open. That is only ~17 percent margin, and TI's thermal-regulation loop (TJ 145 C, 1.25 s timeout) is a second constraint on top. CROSS-BOARD CONSEQUENCE: CAR-REQ-14 puts inrush limiting on the DAUGHTER, so the daughter must limit its own charge current enough that the carrier eFuse never sits in current limit for 162 ms. That number is not currently in ICD-01 and MUST be added before the daughter runs design their inrush circuits. Raise at H2.
P3	ICD-01 re-issued as rev A2: creepage 0.60 -> 0.635 mm board-wide, plus two new binding sections (s6.6 bank-charging contract, s8.4 CAR-REQ-08 realisation)	Coordinator required the changed creepage number to reach the ICD because daughters are already designing to it. Verified the raise does not break the connector: the binding geometry is the 0.84 mm pad-to-pad copper gap on a 1.70 mm annulus, which clears 0.635 mm at 1.32x (was 1.4x vs 0.60 mm). check_creepage.py only demands 0.60 mm so a 0.635 mm layout passes it by construction; 0.635 mm is enforced by the hand-written .kicad_dru rule at P5 (TRAP-1). NEW s6.6 gives the strobe run the numbers it was told to read from the ICD: daughter charge current ≤ 0.25 A (af) / ≤ 0.5 A (at), absolute ceiling 1.0 A and never above 1.0 A for > 162 ms, with the arithmetic showing why (raw 2800 uF bank charges in 134 ms against a 162 ms latch-off timer = 17 percent margin, and the invariant 3.2 J must be dissipated in the daughter's limiter not the carrier's HTSSOP-20). NEW s8.4 states the CAR-REQ-08 realisation and the three review-gate-4 cases explicitly.

Phase	What	Why
P3	RE-TUNE REQUIRED on the eFuse UVLO/OVP string: both proposed thresholds fail at LEGAL PoE rail voltages. Binding requirements set for P4.	Reviewed the part-sourcer's R66/R67/R73 = 1M/12.4k/20k (UVLO 38.2 V rising, OVP 61.9 V rising / 57.9 V falling). TWO functional failures at legal operating points, not margin quibbles. (1) OVP FALLING 57.9 V sits only 0.9 V above the 57 V legal PSE maximum. The TPS16630 OVP reference is 1.2 V rising / 1.122 V falling with +/-2 percent, on top of +/-1 percent resistors - worst case the falling threshold lands near 55.6 V, BELOW 57 V. Consequence: after any overvoltage event a fully legal 57 V rail would never re-enable the daughter. REQUIREMENT: size so the OVP falling threshold stays above 57 V across worst-case tolerance, i.e. nominal falling ≥ 59.5 V, which puts nominal rising near 63.6 V given the fixed 0.935 hysteresis ratio. Upper bound is satisfied: worst-case-high rising ~ 65.8 V is under the 67 V abs max, and the SMBJ58A breaks down at 64.4 V min so the TVS acts first anyway. (2) UVLO RISING 38.2 V is ABOVE the 37 V legal minimum PD input for the af build. A compliant PSE sitting at 37 V would leave the eFuse permanently off and the daughter dead. REQUIREMENT: UVLO rising below 37 V with margin - target ~ 34 -35 V rising, ~ 32 V falling, which still sits above the PD controller's own ~ 30 V drop-out so the two do not fight. P4 pwr sheet sizes both with the datasheet equations and stocked values; the schematic reviewer verifies both against these numbers.

Phase	What	Why
P3	Library adjudication: 4 footprint edits approved (J3 annulus normalise, TPS16630 + SCT2A25 stock swaps, J1 mounting holes non-plated); vias-in-pad deferred to P9 DFM	J3 (2x7 POWER header, the one actually carrying 48 V raw) was pulled with a 1.80 mm annulus giving only 0.74 mm pad-to-pad and 1.165x over the 0.635 mm HV requirement, while J4 (2x12 signal) had 1.70 mm / 0.84 mm / 1.32x - the two frozen ICD connectors were mutually inconsistent and the HV-carrying one had the WORSE margin. Normalising J3 to 1.70 mm restores 1.32x, makes both connectors share one land pattern, and keeps ICD-01 rev A2's published number true, so no ICD re-issue of the gap is needed. Drill is 1.10 mm not the 1.02 mm the ICD assumed, but gap = pitch - annulus so the conclusion is unaffected; ICD text corrected. TPS16630 swapped to KiCad stock HTSSOP-20-1EP with the 3.4x6.5 copper / 2.96x2.96 mask SMD-defined pad because it is the 48 V hot-swap switch and its thermal pad IS its dissipation path - the pulled part had ~24 percent less exposed metal, asymmetric rows and 24-percent-narrow signal pads. SCT2A25 swapped for a proper IPC heel fillet on a 100 V buck's power pins. J1's two dia-3.25 mounting holes were plated with ZERO annular ring and would trip P7 DRC - made non-plated. DEFERRED deliberately: untented vias-in-pad under the ESP32 thermal land (0.075 mm ring, below JLC's ~0.15 mm PTH minimum) and the HTSSOP pad - the P9 DFM gate should adjudicate these against real fab rules with its own fix loop rather than have them silently changed in the library.
P4	D40-D42 PESD3V3L1BA as sourced is BIDIRECTIONAL (UMW/Youtai), not the unidirectional Nexperia part - orientation-free today, but a second-source hazard	Librarian withdrew its own polarity-ambiguity flag after fetching the ACTUAL sourced datasheet: UMW/Youtai, states 'Bidirectional configurations', labels the SOD-323 only 'Pin1 Pin2' with no anode/cathode, no V.F spec, and the package carries no polarity feature or band. The symbol's generic 1/2 pin naming is therefore CORRECT, not a library defect, and either orientation is valid on the ADC0/ADC1/ID_ADC clamps. CARRY-FORWARD RISK: genuine Nexperia PESD3V3L1BA IS unidirectional, so if this line is ever second-sourced to real Nexperia stock the orientation starts mattering and a board built either-way would have one clamp backwards. Must be noted on the BOM line and re-checked at any resource change.

Phase	What	Why
P4	ICD-01 rev A3: par's blocking thermal issue UPHELD and re-derived; sealed enclosure does NOT close for the 802.3at upgrade	The par was right - rev A2's 56 C (af) / 69 C (at) pair was internally inconsistent. Convection is near-linear in delta-T at this scale, so ~2x box heat gives ~2x rise; 56 C implies 3.13 K/W while 69 C implies 2.29 K/W. Re-derived on the par's independently sourced 3.6-4.3 K/W for a sealed non-metallic box (which supersedes my optimistic 3.13 K/W): af ~9.9 W gives 61-68 C at a 25 C room and 76-83 C at 40 C; at ~19.2 W gives 94-108 C at 25 C and 109-123 C at 40 C. CONSEQUENCES: (1) the at upgrade REQUIRES a vented enclosure - 94-108 C exceeds both the ESP32-S3's +85 C ambient limit and the par emitter's +85 C Topr max; (2) even af closes only near a 25 C room and is marginal at 40 C, so the Q13 provisional '0-40 C sealed' is not compatible with EITHER build and must go back to the owner; (3) every internal-air figure now carries an explicit room ambient. Root cause worth noting: the P1 power architect flagged exactly this ('the at upgrade needs vents OR ambient below ~30 C') and it was recorded as a decision but never propagated into the ICD's numbers - a hand-off gap between phases, not a physics error.
P4	CR-1 GRANTED: ID_ADC board-type code space allocated and published in ICD s3.4 (unblocks the par's P4)	Carrier owns the code space per CAR-REQ-07. Carrier fits the 10k top leg to +3V3; daughter fits R_ID to GND; V_ID = 3.3 x R_ID/(R_ID+10k). Codes: strobe = 2.7k (0.702 V), par = 4.7k (1.055 V), plus 7.5k/12k/18k/30k reserved, short = FAULT, open = NO DAUGHTER. Minimum adjacent separation 0.353 V against ~+/-0.05 V worst-case 1% divider error, so a +/-0.15 V detection window is safe. Top code deliberately stops at 2.475 V because the ESP32-S3 ADC saturates near 3.1 V at 12 dB attenuation and anything higher could not be told from open. FAULT and NONE both force ENABLE de-asserted. Also granted CR-3 (optional 14-bit / 4.883 kHz mode) and CR-4 (90 degree phase stagger via LEDC hpoint - free in hardware, avoids a 4x input-current step).

Phase	What	Why
P4	PWM timer contract REWRITTEN (ICD s3.5): the hard PWM0-3/timer0 + PWM4-7/timer1 assignment is withdrawn - it mis-specified the strobe	The strobe's flash gate is a 5-200 ms one-shot, not a duty setting: one LEDC period at 9.766 kHz is 102.4 us, so a 5 ms flash is only ~49 periods. It must drive the line as a GPIO/RMT one-shot or re-program the timer it owns. The old ICD text implied all 8 channels sit at the default frequency, which is false. New contract: the carrier does NOT hard-assign timers; it publishes the hardware ceiling (8 channels, 4 timers, low-speed only, 14-bit max, channels sharing a timer share frequency AND resolution) and each daughter declares its own channel->timer->frequency map, which carrier firmware applies from the ID_ADC code. This also absorbs the strobe's switch to RGBW (now consuming all 8 channels) without another ICD change.
P4	ICD rev A4: sealed-with-wall-conduction enclosure decided; rev A3's 'at needs venting' conclusion WITHDRAWN; thermal now published as a box-air heat BUDGET	Owner decided the enclosure: SEALED, non-metallic, LED heat conducted OUT through the wall (not vented), consistent with H1-Q5. That changes the arithmetic materially. Rev A3 charged essentially the whole PoE budget to box air (~9.9 W af / ~19.2 W at) and unsurprisingly did not close - but the light engine is the large term and in this configuration the LED's ~6 W leaves through the wall, so only in-box ELECTRICAL losses heat the air. Re-derived as a budget rather than a temperature, which is what a daughter can actually design against: holding internal air to 70 C (15 K below the +85 C limit shared by the ESP32-S3 and the par emitters) at 3.6-4.3 K/W gives an allowable box-air heat of 10.5-12.5 W at a 25 C room, 9.3-11.1 W at 30 C, 7.0-8.3 W at 40 C. Carrier spends ~2.4 W (af) / ~3.7 W (at) of that. CONCLUSION: sealed closes for BOTH af and at, but ONLY if the LED heat genuinely leaves through the wall - that is now the load-bearing assumption of the entire thermal case, so ICD s7.7 requires every daughter to DECLARE its box-air vs wall split, and states that the wall-conduction path is a mechanical requirement rather than a nicety. If it is not built, the numbers revert to rev A3 and at does not close.

Phase	What	Why
P4	CR-5 GRANTED and recorded as firmware requirement FW-01 in new ICD s7.8; four further firmware requirements of record captured alongside it	Owner committed carrier firmware to PWM-domain dithering of ≥ 3 -4 bits - it is the mechanism by which the par's PAR-REQ-01 is met and no hardware on that board can close it, so it is a requirement, not an option. Recorded in ICD s7.8 with the critical qualifier that it MUST be PWM-domain and NOT 60 fps frame-domain, because a 4.4 percent dither at 60 Hz breaches IEEE 1789's no-effect level by itself and would create the very flicker it is meant to avoid. Captured four more firmware commitments in the same table so none can be silently dropped when firmware is written: FW-02 = 90 degree hpoint stagger for 4+ channels on one timer (CR-4); FW-03 = read ID_ADC at boot and never assert ENABLE on FAULT or NONE; FW-04 = apply the daughter's declared timer map rather than assuming the default rate; FW-05 = maintain MPS ≥ 10 mA DC, which forbids ESP32-S3 deep sleep while PoE-powered. These will also be named in the DOC-01 design document.
P4	CORRECTION to my own record: the rev A2 56 C af figure was NOT wrong - it was correct under an unstated assumption. Only the at figure was a genuine error. ICD rev A5.	I characterised both rev A2 thermal figures as wrong at rev A3. That was unfair to the af figure and is now corrected. The strobe run applied the par's measured sealed-box resistance to its own heat sources and reproduced BOTH disputed numbers from a single binary - whether LED heat leaves through the wall or stays in the box. LED-in-box at 8.5 W gives 56-62 C at a 25 C room, which IS the rev A2 af figure to the degree. So that number was never miscalculated; its two assumptions (LED heat inside the box, 25 C room) were simply never written down, and that omission is the entire defect - it made a correct figure look arbitrary while silently assuming the arrangement the owner has since ruled out. The at figure remains a genuine error: 69 C follows from neither case, since 56 C implies 3.13 K/W and 69 C implies 2.29 K/W. The par's independent 89-115 C is retained as the CONSERVATIVE BOUND and explicitly not averaged with anything.

Phase	What	Why
P4	OPEN ITEM D CLOSED: the 802.3at upgrade does NOT require enclosure ventilation. Sealed closes to a 40 C room.	With the wall-conduction path working, the strobe measures only ~1.894 W of daughter heat into box air, not the ~8.5 W my earlier analysis charged. Recomputed carrier total box-air heat: ~2.4 W overhead + ~1.9 W daughter = ~4.3 W (af); ~3.7 W + ~1.9 W = ~5.6 W (at). Against the allowable budget of 7.0-8.3 W at a 40 C room, the at case reaches 40 + 5.6 x 4.3 = 64 C internal air WORST CASE - under the 70 C design target and well under the +85 C limit shared by the ESP32-S3 module and the par emitters. So the rev A3 'at requires venting' conclusion is withdrawn and open item D is closed. This is entirely contingent on the LED thermal path to the wall actually being built: if it is not, the numbers revert to the LED-in-box row and the margin disappears, so ICD s7.7.3 states the wall path as a MECHANICAL REQUIREMENT and requires every daughter to declare its box-air vs wall split. Also resolves CAR-REQ-18 (two heat sources, one enclosure) favourably - the second source is 1.9 W, not 8.5 W.
P4	P4 reviewer triage: 1 error to fix-loop, 7 fixed in copper, 2 fixed by ICD amendment (rev A6), 7 waived, 2 raised for the owner at H2	Fresh-context reviewer returned 19 findings and independently verified all ten known-critical items (R69 fail-safe, TPS2378 EP-to-VSS, APD-to-RTN, the 1.2V divider, EN abs-max chain, W5500 RSVD asymmetry, GPIO45, magjack map, both frozen ICD pin maps) as correct. MOST IMPORTANT FINDING is F7 and it is a HONESTY item, not a defect: because lib_pin_types types every non-supply pin 'passive', KiCad's output-vs-output, input-not-driven and driver-conflict checks are structurally INERT - the ERC 0/0 proves only that no power.in lacks a driver and no wire dangles. Real pin-level coverage came from the reviewer's per-IC datasheet audit, NOT from ERC, and the design document must not claim otherwise. F6 and F3 were fixed by amending the ICD rather than the board: publishing FAULT's real ~4.3 mA sink requirement (vs the 0.33 mA implied) is cheaper for daughters than a carrier respin plus a GPIO, and the IMON governor claim needed softening because the pin's specified accuracy starts at 0.6 A while the rail's sustained limits are 0.25/0.5 A - both below the floor.

Phase	What	Why
P4	OVP divider upgraded to Yageo RT 0.1 percent 25ppm/degC THIN-FILM (R66/R67/R73). Tempco was the real problem, not initial tolerance.	Owner took the 0.1 percent recommendation at H2. Selecting thin-film rather than merely tighter thick-film is the part that actually fixes it: the reviewer's failure mode was 100 ppm/degC over a 60 degC excursion dragging the OVP FALLING threshold to ~56.5 V, BELOW the 57 V legal PSE maximum, so a legal rail could fail to re-enable after an overvoltage event. Yageo RT is 0.1 percent AND 25 ppm/degC, a 4x tempco improvement, which moves worst-case falling from ~56.5 V to ~58.2 V - roughly 1.2 V of real margin instead of 0.18 V. C865592 620k / C110775 10k / C865172 12k, all 0805, stock 5453 / 811444 / 25480. R4/R5 (T2P level shift on the poe sheet) deliberately stay on the 1 percent C17414 line - only the threshold-setting string needs precision, and spending 0.05 USD each on level-shift pull-ups would be waste.
P4	BOM resync tooling hardened: bom_sync.py now preserves not_fitted lines	The first sync correctly rebuilt refs/qty from the netlist but DROPPED C334927, the 63.4R class-4 upgrade resistor, because it has zero refs by design - it is deliberately not fitted. That would have silently deleted the single component that makes D-01's 'resistor change, no respin' promise real. Lines now carry a not_fitted flag and survive the sync. Final state is stable and self-consistent: 50 BOM lines, 109 placed components, 0 lines changed on a repeat run, 0 netlist LCSC codes without a BOM line, and the only component without an LCSC property is H5 (a mounting hole, board-only by design).

Phase	What	Why
P5	BLOCKED before P5: the mag-jack qualification found that NO available 10/100 PoE+ part fully meets the goal the swap was ordered to achieve	Best documented 10/100 PoE+ ICM is Wurth 7499410213: publishes '600 mA per centre tap, compliant with IEEE 802.3at' (datasheet rev 002.000 p3), -40..+85 C. But 600 mA is EXACTLY the 802.3at DC maximum and BELOW the 0.686 A peak, and the agent found no 10/100 part publishing more - everything higher is gigabit or PoE-bt. So the swap does not fully deliver D-01's at-upgrade promise either. It also is NOT a drop-in: no integrated bridge (needs two external bridges), 12 of 14 pins move (P1/P2 RX->TX, P4/P5 TX->RX, P7/P8 no-connect->live Mode A taps, P9/P10 rectified->raw spare pairs), an internal Bob Smith network that cannot be removed and contradicts poe.py's rationale, and the chip-side/line-side isolation pad gap collapses from 3.58 mm to 1.05 mm (below HALO's 1.40 mm guidance, and check_creepage will NOT catch it because it only enforces 0.635 mm for a 57 V working voltage). Cost +7.20 USD/board = +103 USD on 14, moving the carrier from ~26-32 to ~33-39 against a 30 USD target. LCSC has it at ZERO stock so JLC cannot source it: J1 becomes DNP, ordered from Digi-Key, hand-soldered 16 terminals x 14 boards. WORST PART: at 600 mA the magjack becomes the LOWEST-rated element in the 48 V path, and neither the eFuse (1.0 A) nor the PD hot-swap (~1.0 A) protects it - it can sit at 1.0 A for 162 ms before latch-off. Protecting it by lowering R(ILIM) at the at build makes D-01 a TWO-resistor upgrade, which is precisely what D-01 exists to avoid.

Phase	What	Why
P4	Gigabit screen SUCCEEDED: LINK-PP LPJG0926HENL (C22457393) publishes 720 mA @57VDC continuous - the first candidate that covers the 0.686 A at PEAK, not just the DC max - and JLC stocks and places it	The earlier rejection of gigabit ICMs was a category error, and re-running it changed the answer. LINK-PP LPJG0926HENL publishes 'DC Current/Voltage Rating pse Pins: 720mA MAX @57VDC (Continuous)' on p1 of drawing LP18022610 rev A. That is 1.20x the 802.3at 0.600 A DC max and 1.05x the 0.686 A peak. LCSC C22457393, 3109 in stock, 3.7852 USD at the qty-10 break = 52.99 USD for 14, and critically JLCPCB Assembly Type = Wave Soldering, so JLC PLACES it - no DNP, no hand-solder, single purchase order. That is 80.57 USD CHEAPER than the Würth plan while RAISING the rating and REMOVING the partial-population problem. It also largely rescues D-01: because 720 mA exceeds the at peak, there is no need to lower R(ILIM) for normal at operation, so the one-resistor upgrade promise survives (only bounded transient fault overstress at the eFuse's 1.0 A limit remains). Thermal cross-check against ICD s7.7: its 0..+70 C rating gives +15.9 K margin at the 30 C room ceiling the ICD already sets for at, and +20.9 K at 25 C. Costs carried: no integrated bridge (2 external bridges - unchanged from the Würth plan), isolation gap 1.06 mm NOT fixed but mitigable to 1.36-1.56 mm because the two pads nearest the 48 V nodes are TD4+/TD4- chip-side windings of a pair the W5500 never uses, and de-specs on return loss (-16 vs -18 dB), CMRR (-30 vs -35 dB) and op temp.

Phase	What	Why
P4	DECIDED (owner delegated): adopt LINK-PP LPJG0926HENL (C22457393) as J1. Closes OPEN-A.	Owner delegated all remaining decisions. Option (a) dominates (b) Wurth on every axis: 720 mA published vs 600 mA, covers the 0.686 A at PEAK vs falling under it, 52.99 vs 133.56 USD for 14, JLC wave-solders it vs DNP+hand-solder x14, 3109 LCSC stock vs zero. It beats (c) keep-HY because it publishes a rating at all, and it rescues D-01's one-resistor promise (720 mA > 0.686 A peak means R(ILIM) need not drop at the at build). Accepted costs, all recorded: no integrated bridge so two external bridges are needed (same as the Wurth plan, not new); isolation gap 1.06 mm is NOT fixed and MUST be mitigated by leaving pins 9/10 (TD4+/TD4-, chip-side windings of a pair the W5500 never uses) as bare no-connect pads shrunk to 1.30 mm, giving 1.36-1.56 mm to any energised net - this must be re-checked at P8 and check_creepage will NOT catch a regression (LEARNINGS 2026-07-28); de-specs on return loss (-16 vs -18 dB) and CMRR (-30 vs -35 dB), acceptable at 100BASE-TX on 2 of 4 pairs; op temp 0..+70 C makes it the lowest-temperature-rated part on the board, with +15.9 K margin at the 30 C room ceiling ICD s7.7 sets for at. Residual risk accepted: LINK-PP's wording is 'pse Pins ... 720mA MAX' rather than 'per centre tap'; the part has exactly four power pins each being one transformer's line-side CT, so per-pin is the only sensible reading.

Phase	What	Why
P4	Magjack swap COMPLETE and verified. Two library defects fixed beyond the brief; my own brief contained a wrong pin table and the agent correctly overrode it.	PROCESS NOTE FIRST: my Task-4 pin table was WRONG - I carried the Wurth 10/100 part's pin deltas into a brief for the LINK-PP, which is a different part with a different pin count. The agent detected the contradiction, built to the datasheet (VC1-4 = pins 11-14, LEDs 15-18, EH 19-20 - confirmed against parts/C22457393.json) and reported the override rather than splitting the difference. That is exactly the behaviour that keeps a wrong magjack pinout from flipping the connector end-for-end. FIXED BEYOND THE BRIEF: (1) board-lock <-> VC creepage was 0.487 mm against the board's own 0.635 mm rule - this one IS voltage-derived so check_creepage would have failed it at P8 after routing. Fixed by making pads 19/20 oval 2.00x2.60 (narrow toward the taps, tall for insertion force) and VC1/VC4 1.20 mm: gap now 0.687/0.697 mm with every annular ring at the 0.15 mm JLC minimum. The chip-side/line-side barrier IMPROVED as a side effect, 1.401 -> 1.451 mm, now clearing HALO's 1.40 mm guidance. (2) R3's land moved 0603 -> 0805 (C3000584 fitted, C334927 not-fitted alternate) so the 63.4R class-4 upgrade at ~105 mW is in-spec on the SAME land - D-01 promises a resistor swap with no respin and an out-of-spec alternate would not deliver it. Bridges: 2x ABS210 (C2892567), selected on RthJA 65 C/W giving Tj 118-133 C; MB6S rejected by measurement at 90 C/W and Tj 176 C. RDEN re-centred 24.8k -> 24.2k now that the bridge diode I-V is known. ERC 0/0, netlist_audit exit 0, BOM 52 lines / 111 placed / 0 orphans.

Phase	What	Why
P5	P5 complete: 100x80 mm, 4-layer JLC04161H-3313, 3 mm corner radius, 4x M3. Three library defects found and fixed that placement could never have resolved.	board_init report read back rather than assumed: corner_radius 3.0 honoured (NOT clamped to the mounting-hole inset), mounting_holes 4, outline_bbox [19.58, 57.132, 119.58, 137.132] = exactly 100.0 x 80.0 mm. Setup violations fell 92 -> 36 after fixing: (1) ESP32 thermal-land vias-in-pad at 0.075 mm ring / 0.25 mm drill, failing BOTH min-annular and min-hole. My first attempt enlarged them to 0.60/0.30, which fixed those two but created 24 clearance + 24 mask-bridge errors instead because the via pads then touched the thermal sub-pads - so I REMOVED them entirely. Thermal vias under a ground land belong to the board, not the footprint: P7 stitching places them against the real GND pour with correct net assignment, and it also removes the via-in-pad solder-wicking risk that was already on the P9 list. (2) D1 SMBJ58A courtyard was 'malformed (self-intersecting)' from zero-length fp_line segments - an easyeda2kicad artifact; scanned board-wide and removed 5 across 2 footprints. All 36 remaining violations are components stacked at the origin pre-placement (27 copper_edge_clearance, 3 courtyards_overlap, 3 mask bridges, 1 npth-in-courtyard) plus 2 benign warnings. ZERO library defects remain. Also verified by hand (fp_verify does not model drills): J3/J4 both 1.700 mm annulus / 1.100 mm drill / 0.840 mm gap / 1.32x / 0.300 mm ring, matching ICD s5.2 exactly. Library completeness re-derived from the filesystem after the lib_pull bug: 112/112 components resolve a footprint.

Phase	What	Why
P6	P6 place gate PASS. My own DRU rule was over-scoped and I fixed it; a blanket silk sweep was over-aggressive and I caught it before committing.	Place gate 0 violations; hpwl 4212 -> 3601 mm, crossings 955 -> 613, utilisation 24.7 percent; route probe completion 0.9842 (5 unrouted of 317), above the 0.98 bar. TWO SELF-INFLICTED ISSUES FIXED: (1) my hand-written .kicad_dru HV rules had no pad-pair exclusion, so they fired 30 times between ADJACENT PINS OF THE SAME PACKAGE - U1's own VDD/VSS at 0.20 mm, U22's HTSSOP pitch, C1/C61/C62/C63's own two pads. No placement can fix a 0.65 mm-pitch package, and that was never the rule's purpose: TI's 0.635 mm guidance governs BOARD copper around the part, not the vendor's lead frame. Added !(A.Type=='Pad' && B.Type=='Pad') to all three rules - 30 errors to 0 - and this loses no coverage because P8 check.creepage independently models pad-to-pad from the voltages table. (2) My first silk-clip pass was scoped board-wide and removed 241 primitives, stripping legitimate silkscreen from nearly every passive, because a 0603's body outline naturally overlaps its own pad bboxes. Caught it by sanity-checking the count before committing, restored from backup, and re-scoped to only the refs DRC actually flagged: exactly 5 clipped (J1 x4, D22 x1), verified by differencing silk-layer references against the backup. Board DRC is now 0 non-unconnected violations. Also frozen: ICD s7.2 connector coordinates (J3 origin 23.00,76.00; J4 origin 71.00,76.00; H5 46.00,74.00), and J1 moved 5.75 mm left onto the ICD nominal envelope, taking the daughters' notch margin from 0.13 mm - not a margin - to a real one.

Phase	What	Why
P7	MEASUREMENT RECORD SETTLED (shape-aware). My P4 figure was RIGHT; my P7 'correction' was WRONG. J1 barrier = 1.613 mm, passes HALO. Shield-to-tap = 0.603 mm, passes IPC but 0.032 mm short of my adopted 0.635 mm.	I produced three numbers for the same barrier because I twice used a formula that did not match the pad SHAPE. J1's signal and tap pads are CIRCLES, and for circle-circle the radial form $\text{hypot}(dx,dy) \cdot (r1+r2)$ is EXACT - so my original P4 figure of 1.451 mm was correct all along and HALO's 1.40 mm was met. My P7 'correction' to 1.148 mm applied a RECTANGLE model to circular pads, which underestimates diagonals, and I retracted a correct number and reported a false failure. The router caught it. Definitive figures from a single shape-aware tool (work/p7/pad_gap.py, capsule model: circle = degenerate segment, oval = stadium spine, rect = per-axis): BARRIER pad 2 (/ETH.TXN) to pad 11 (/poe/POE_TAP_A1) = 1.6130 mm as built, or 1.4510 mm with pad 2 at its original 1.524 mm - BOTH pass the 1.30 mm rule and HALO's 1.40 mm. SHIELD board-lock 19 (oval) to tap 11 (circle) = 0.6029 mm, and 20 to 14 = 0.6127 mm - not the 0.550 mm I reported (rectangle model on an oval-circle pair) and not the 0.687 mm of P4 (radial model on the same). CONSEQUENCE FOR THE RECORD: 0.6029 mm PASSES IPC-2221B B2 (0.60 mm for the 51-100 V band) and falls 0.032 mm short only of the 0.635 mm I adopted board-wide from TI's PD-specific VSS-to-VDD guidance - which is about copper around the PD controller, not a magjack shield tab. That is a far smaller deviation than I reported and it clears the actual standard.
P7	DELIBERATE NON-REVERSION: J1 pad 2 stays at 1.200 mm even though the shrink that produced it was based on a wrong model	The shrink from 1.524 to 1.200 mm was unnecessary - the barrier met HALO's 1.40 mm at the original size (1.4510 mm). Reverting would restore 0.162 mm of annular ring per side on a THT MDI pad. I am NOT reverting, for three reasons: (1) 0.150 mm ring on a 0.900 mm drill is exactly JLC's stated PTH minimum, so the as-built pad is within spec, not out of it; (2) the barrier is BETTER as built, 1.6130 vs 1.4510 mm; (3) enlarging the pad now would shrink every pad-to-track gap around it on a board the router just brought to 0 clearance errors with a measured 1.4221 mm worst case, forcing a re-verification pass for margin that is not needed. Recorded so the next reader knows the pad size is a deliberate accepted state and not an unexamined leftover.

Phase	What	Why
P7	C35 moved +0.300 mm in y (85.513 -> 85.813): U10 GND fan-out corridor 0.681 -> 1.681 mm. Placement-blocked routing unblocked without touching the planes.	The router quantified the block precisely: U10's south pad row bottom at y 83.6819 vs C35 pad-1 top at 84.3627 left a 0.6808 mm corridor, while the smallest DRC-legal via here (min_via_diameter 0.5) needs $0.5 + 2 \times 0.2 = 0.900$ mm - short by 0.192 mm. It recommended C35 +0.5 mm; that overlapped U30's courtyard by 0.18 mm ² and failed the place gate, so I used the gate as the oracle and backed off to +0.300 mm, the largest offset that passes. Corridor is now 1.6811 mm, nearly 2x what a legal via needs. Verified with BOTH oracles: place gate PASS (0 failing), check_decoupling unchanged at the 2 pre-existing C1 warnings with no new ones - which matters because C34/C35 exist to sit close to U10's supply pins, so moving one could have broken the thing it is for. This is the cheap lever the router flagged as preferable to permitting inner-layer routing, and it costs no plane integrity.
P7	MY C35 MOVE WAS WRONG AND INTRODUCED A SHORT. Both oracles I validated against passed it. Router caught and reverted it.	Two compounding errors. (1) MEASUREMENT: my 0.681 -> 1.681 mm corridor figure read C35's FOOTPRINT origin y, not its PAD-1 CENTRE y - off by exactly 0.700 mm. The real corridor was 0.9811 mm, so the move bought far less than I claimed. (2) WRONG ORACLE, and this is the serious one: at y=85.813 C35's pad 1 landed on an existing /ETH_RSTn run at y 85.55, giving +6 DRC errors including a SHORTING.ITEMS and a solder_mask_bridge. I validated with 'gate place' (PASS) and check_decoupling (unchanged) - and BOTH passed, because placement-phase checkers do not see routed copper. On a board that is already routed, the authoritative oracle for a placement change is DRC, not the place gate. I moved a part on 3089 segments of existing routing and checked it with tools that are blind to tracks. (3) The corridor could never have closed the GND item anyway: the real blocker was a VOID in the In1 GND fill created by neighbouring +3V3 plane vias' clearance holes, so a via there would have connected to nothing - 0 legal 0.5 mm centres across the whole corridor at 0.025 mm step. The router solved it properly by escaping U10 pad 9 NORTH into the LQFP's own die shadow. This is my fourth verification error this phase and the only one that damaged the board; every one was caught by an independent check rather than by me.

Phase	What	Why
P8	SILK LEGIBILITY requirement accepted and QUEUED (not yet done): refdes must sit within ~1-1.5 mm of their own part while silk overlap/over-copper/edge stay at zero. Cause confirmed as the P6 sweep.	CAUSE CONFIRMED FROM MY OWN HISTORY, not taken on trust: reports/place/silk*.ops.json is 9 files carrying 160 move_text ops, which the P6 placement agent used to take silk violations 236 -> 5. It pushed refdes outward to clear overlaps and nothing pulled them back, so the board was optimised for a zero-warning gate at the cost of assembly legibility. BASELINE (coordinator-measured, my script agrees on the worst offenders): 116 footprints, 116 visible refdes, median offset 4.079 mm, max 12.955 mm, 100 of 116 refdes more than 1 mm beyond their own part's extent. Worst: C2 12.767 mm offset on a 2.699 mm part = +10.067 mm excess; C62 +8.168; R101 +7.551; C81 +7.377; C1 +6.666. A 0603 whose label sits 8.6 mm away is closer to three other parts than its own and is unattributable on a populated board. WORK NOT STARTED, DELIBERATELY: the P8 verify fixer is concurrently writing the same .kicad_pcb (copper: 191 check_current, 19 check_creepage, 8 check_return_path). move_text and copper edits both rewrite the file, so running them concurrently risks a torn write - the same class of failure as the P4 librarian race. Silk work must SERIALIZE after the fixer completes. Approach when it runs: constrained minimisation, not a sweep - batches of move_text with DRC after each, revert any batch that regresses, then re-run drc_routed and confirm 0/0 before verify. Where a label cannot be placed unambiguously, prefer the position whose nearest OTHER part is clearly further away, because a label adjacent to the wrong part is worse than a distant one. Measurement tool at work/refdes.prox.py is the before/after objective.

Phase	What	Why
P8	NEW SAFETY DEFECT the gate CANNOT see: POE_TAP_A1 to POE_TAP_A2 = 0.3295 mm on all four layers, with full line voltage between them. Third instance of the same class.	check_creepage computes $dv = v_{hv} - v_{other}$ and skips any pair with $ dv \leq 30$ V. Both A- side taps are declared 57 V, so $dv = 0$ and the pair is NEVER CHECKED - yet A1 and A2 are the two AC INPUTS of bridge D2, so the full line voltage sits across them. Measured 0.3295 mm on all four layers (the two transition vias at (49.650,74.213) and (49.200,73.400), 0.9295 mm centre-to-centre with 0.6 mm pads) against 0.60 mm outer / 0.20 mm inner required. B1- B2 measures 1.4778 mm and is fine, so it is one via pair. THIS IS THE THIRD TIME on this board that a safety check reported nothing be- cause it was looking at the wrong thing: (1) my DRU rule named V48_RAW when the barrier lives on POE_TAP_*; (2) constraints.json volt- ages omitted the tap nets entirely so the shield- to-tap gap was invisible; (3) now the voltages MODEL cannot express a differential pair, so two nets at the same declared potential with mains-class AC between them are skipped by construction. Two fixes, both owner decisions: declare each bridge's negative-side tap at -57 V (as V48_RTN already is) so the checker sees a 114 V pair - this adds gate failures rather than hiding them, which is the honest direction; or move the A2 via ~0.30 mm NW to fix the geom- etry. NOT ACTIONED - I am at my context limit and will not start a copper change I can- not verify.

Phase	What	Why
P8	TRIAGE REOPENED on two P8 residual clusters my predecessor closed by classification rather than by attempt: 5 of the 15 check_creepage errors are NOT package-internal, and 5 of the 53 undersized tracks carry 1.0-2.0 A on 0.20 mm copper	Read work/p8/creep_located.json item-by-item instead of trusting the summary. Ten of the fifteen creepage errors are genuinely unfixable pad-to-pad inside one package (U1 SOIC-8 VSS pad 9 against pads 1,2,3,5,6,7,8 all at exactly 0.200 mm = the lead pitch; U20 0.295 mm; U22 0.375 mm and 0.250 mm). But FOUR are PAD-to-TRACK and ONE is TRACK-to-TRACK, all on the /poe/SHIELD or /poe/LED_Y_A nets inside J1 at 0.203-0.220 mm. A track can move; a lead frame cannot. And the connector's own pad-to-pad shield-to-tap gap is 0.6029 mm measured shape-aware, so J1 itself holds 0.60 mm - it is the ROUTED shield copper that erodes it to 0.21 mm. Second cluster: residual_report.json marks all 53 undersized_track items blocked, which is true for widening IN PLACE (the room figure is the widest each segment can be), but the blocker in five cases is a 0.2 mm signal net or a switch node that can move instead. +12V carries 2.0 A through a 0.20 mm x 1.98 mm segment at (66.35,112.05)-(67.75,113.45) in series with the J3 feed - 0.20 mm of 1 oz copper is worth about 0.3 A at dT=10 C, so that is a 6x overload and the worst single number on the board, not a modelling artifact. Dispatching both as ONE bounded fixer order with a 2-attempt-per-item budget and an explicit instruction that a measured unfixable is a better answer than a reroute that trades EMI for a width figure.

Phase	What	Why
P8	A-SIDE TAP CREEPAGE FIXED: POE_TAP_A1 to A2 = 0.3292 -> 0.6502 mm (F.Cu) / 0.6524 mm (In1, In2, B.Cu) against 0.60 mm required. drc_routed still 0/0, verify failing multiset BIT-IDENTICAL to baseline. A SECOND hidden pair on the same two nets was found and fixed in the same pass.	Two vias moved and dropped 0.6 -> 0.5 mm (min_via_diameter 0.5, annular 0.1 mm = board and JLC 4-layer minimum; via 40e3f5b2 on this same net was already 0.5/0.3, so this is existing practice). The pocket is boxed in on every side - SHIELD board-lock oval J1-19 to the NW, SHIELD B.Cu escape to the NE/E, and the DRC-enforced magjack_isolation_barrier 1.30 mm rule to the S - so the fixer grid-searched both positions maximising the WORST margin rather than the headline gap. Its first optimum gave a 0.726 mm tap gap but left only +0.005 mm to SHIELD, i.e. it robbed one 57 V gap to pay another; the balanced answer equalises them at ~0.65 mm. Final margins: tap +0.05 north / +0.13 south, SHIELD +0.047, ETH barrier +0.060. The 1.30 mm centre-to-centre target in my work order is UNREACHABLE in the north pocket (achieved 1.136 mm) and I was wrong to ask for it. MOST IMPORTANT: the work order named ONE pair and there were TWO. A1's south transition via 17a75fde sat 0.5500 mm from A2's B.Cu run and 0.5826 mm from its corner - also a differential 57 V pair, also below 0.60 mm, and invisible until the 0.3292 mm pair was fixed because a worst-pair- only sweep hides its siblings. Fixed in a sec- ond separately-DRC-verified edit. B-side con- firmed unchanged at 1.4778 mm on all four lay- ers. That makes FOUR distinct instances on this board of a safety check reporting nothing because it was aimed at the wrong thing, and the fourth was hidden by the third.

Phase	What	Why
P8	REOPENED TRIAGE PAID OFF: check_creepage 15 -> 11 and check_current under- sized_track 53 -> 51, drc_routed held at 0/0 throughout. Four of the five 'package-internal' creep- age errors were routed SHIELD copper and were fixable; two 0.200 mm power-trunk segments now carry 0.500 mm.	The four J1 SHIELD-to-tap items were NOT package geometry. Their gaps went 0.213 -> 0.6390 (A1), 0.2142 -> 0.6388 (A2), 0.220 -> 1.7700 (B1), 0.2129 -> 0.6620 (B2) mm by re- placing the whole 39-segment SHIELD pad19- ->pad20 staircase with a 6-segment polyline at the analytic optimum. The corridor is genu- inely tight - the pad11/pad12 axis is 2.8398 mm and needs $1.462 + 1.300 = 2.762$ mm, leaving 0.0778 mm of TOTAL slack - and the topology was forced rather than chosen (pads 12/13 escape south on F.Cu; the +3V3 via at (45.800,73.450) blocks the north corridor). Re- swept afterwards: 0 of 31 SHIELD items now under 0.600 mm, and the remaining net-pair minima are the connector's own land (0.6029 / 0.6127 mm) which no layout can change. B3 and B4 went 0.200 -> 0.500 mm identically: a signal net hops to B.Cu across the V48_RAW trunk with its two vias 1.050 mm off centre- line, and pushing each 0.200 mm out restores 0.700 mm, i.e. +0.065 mm on the 0.635 rule. All five cluster-B segments were confirmed to be BRIDGES in their net's connectivity graph, so each really does carry the full rail current - no parallel path was assumed. Two RE- JECTIONS worth recording because both were the right call: the fixer's first attempt at the A2/LED_Y_A item did reach the target but only by pulling A2 to 0.2259 mm from J1's board lock, down from 0.7065 mm - trading one 57 V clearance for another - and it rejected its own result after measuring that every other 57 V tap holds ≥ 0.6608 mm there. And B1/B2 CAN reach 1.1000 mm exactly, but only by spreading the 3.3 V buck switch-node crossing from 1.2024 to 2.100 mm, which is the EMI-for- width trade the work order forbade.

Phase	What	Why
P8	ITEM 5 (POE_TAP_A2 to LED_Y_A, 217 pairs at 0.2031 mm vs 0.600 mm required) LEFT OPEN AND UNRESOLVED, not waived and not closed. Carried to H4 as an owner decision. Prioritised the silk legibility pass ahead of a third solve attempt.	Three things had to be separated. (1) The authorisation to trade A2's board-lock margin was VOIDED by the coordinating orchestrator after the fixer applied, tested and RETRACTED its own result: the +0.0037 mm price came from an ad-hoc seg-seg distance primitive that returns a positive number for INTERSECTING segments, so the trade was never available at that price. Nothing of item 5 is on the board; drc_routed is back to 0/0 and check_creepage is 11. The retraction is the right outcome and the fixer surfacing it unprompted is the behaviour that matters. (2) The 9.8 micrometre impossibility proof is VALID but NARROWER than first stated: it proves only the 0.6608 mm floor-enforced variant impossible, as a 1-D packing bound independent of the primitive bug - every term is a fixed pad radius, a fixed rule or an unneckable PoE net, saturating at -0.0223 mm. The 0.600 mm variant is UNRESOLVED, not impossible: a crossing-aware 3-net solve (A2 + LED_Y_A necked + LED_G_A with the ordering LED_G_A LED_Y_A A2 explicit at every x) has never been run. (3) SCOPE CHECK ON THAT SOLVE, worth recording before anyone runs it: the proposed lever is necking LED_Y_A to 0.100 mm, which is BELOW JLC's 0.1016 mm minimum trace width - the same limit that already generates 189 DFM errors on this board - so the solve must be run at ≥ 0.1016 mm. The 1.6 micrometre correction is immaterial against the claimed +0.4165 mm of room from adding LED_G_A, so the solve remains plausible; it is simply not free. DECISION: the silk legibility pass is an explicit owner requirement, measured, and still entirely unstated, while item 5 has already consumed two bounded attempts. Silk goes first. If budget remains after silk, verify and the reviewer, the 3-net solve is the next thing to spend it on. Presented at H4 as a disclosed shortfall with the 217-pair count and the 0.2031 mm worst case, because a creepage requirement missed by 3x on a 57 V net is the owner's call, not mine. Note for that decision: the J1 board locks have NO declared potential - netless, so no gate models them - which is why the competing margin is a judgment call at all.

Phase	What	Why
P8	SILK LEGIBILITY DONE: 111 in-scope refdes pulled in. Median offset 4.079 -> 1.600 mm; refdes more than 1 mm beyond their own part extent 95 -> 3 (excluding out-of-scope H1-H5); more than 1.5 mm beyond, 86 -> 1. 116/116 still visible, drc_routed 0/0 after every one of 5 batches.	The agent did NOT trust my normalizer targets and was right not to: applied verbatim they produce 283 DRC warnings implicating 85 of 111 refs, because lib_refdes_norm.py uses size/2 for the text half-height and ignores each footprint's OWN silk outline. It rebuilt the geometry from KiCad itself - a read-only bundled-python probe emitting board-frame pad and silk polygons plus each label's true INKED stroke box via TransformTextToPolySet - then ran a greedy solver minimising distance to the label's own footprint origin subject to no contact with any silk, pad mask aperture, other label or the board edge, with a HARD filter that the label's own part must be the nearest part. THE LOAD-BEARING MEASUREMENT: at size 1.0 / thickness 0.15, GetTextBox() returns 1.6965 mm of height but DRC's inked stroke box is 1.162 mm - so using GetTextBox over-constrains by about 0.27 mm per side and makes correct targets look infeasible. That single fact is the difference between this working and not. Every candidate set was validated in a throwaway copy of the board (pcb+pro+dru+sch reproduces drc_routed exactly at 3.4 s per run) before the live board was touched, which is what made 'DRC is the only oracle' affordable across 111 labels. Verified non-regression by snapshot diff: 3294 tracks, 359 vias, 8 zones, 8 board graphics, every pad, every footprint silk and every value field byte-identical; 0 footprints moved, 0 refdes hidden or relayed, no mirrored_text warnings. Attribution beat closeness on 17 refs, 110 of 111 have their own part as nearest, and 5 previously upside-down labels were corrected to board 0/90 deg. ONE RESIDUAL above 1.5 mm: C35 at 3.861 mm, boxed in on four sides (U10's LQFP silk hull 0.32 mm above, U30's module hull 0.45 mm below, C37's hull 0.79 mm east) and its only wide-enough channel is also C34's ONLY attributable slot - C34's alternative escape is over 5.5 mm and equally unattributable, so C34 keeps it. No position within 7 mm has C35 as nearest part. That is a P6 placement consequence, not a silk problem, and per scope it was not relieved by moving a part.

Sheet plan

LUM-CAR-A - hierarchical sheet plan

Five child sheets under a thin stitching root. ~110 components across five clearly separable domains, so hierarchy buys parallel P4 sheet generation and keeps the PD front end's 48 V domain visually and electrically isolated from the logic.

The sheet NAMES below are contractual. They appear verbatim in netlist net names

(/pwr/SW) and `constraints.json` already references sheet-crossing names; renaming a sheet re-names nets and breaks `netlist_audit --constraints`.

Sheet	File	Blocks	Interface nets (sheet pins)	pwr.base
root	'lumina-carrier'	'lumina-carrier.kicad_sch'	none (stitching only)	- 1 (unused)
'poe'	'poe.kicad_sch'	RJ45 PoE magjack, TVS, PD interface, detect/class network, CBULK, T2P level shift, shield hybrid	'ETH_TXP', 'ETH_TXN', 'ETH_RXP', 'ETH_RXN', 'ETH_LED_LINK', 'ETH_LED_ACT', 'T2P'	100
'eth'	'eth.kicad_sch'	W5500, 25 MHz crystal group, EXRES1/TOCAP/1V20, MDI TVS array, SPI series R	'ETH_TXP', 'ETH_TXN', 'ETH_RXP', 'ETH_RXN', 'ETH_LED_LINK', 'ETH_LED_ACT', 'ETH_SCLK', 'ETH_MOSI', 'ETH_MISO', 'ETH_CS', 'ETH_INTn', 'ETH_RSTn'	200
'pwr'	'pwr.kicad_sch'	48->12 buck, 12->3.3 buck, 48 V eFuse, bleed, PG/FLT LEDs	'ENABLE', 'FAULT', 'IMON'	300
'mcu'	'mcu.kicad_sch'	ESP32-S3-WROOM-1-N8, decoupling, EN/BOOT network, recovery header, status LED	everything below	400
'expansion'	'expansion.kicad_sch'	J3 power connector, J4 signal connector, I2C/FAULT pull-ups, ID divider, ADC/ID clamps, H5	'PWM0..7', 'DSPI_*', 'I2C_SCL', 'I2C_SDA', 'ADC0', 'ADC1', 'ID_ADC', 'ENABLE', 'FAULT'	500

1. Net naming - the CANONICAL final netlist names

Mechanism (verified in this repo: `tests/s7_regen/hierdemo/kicad/hierdemo.net`):

1. **Power SYMBOLS make a net global across the whole hierarchy** with a BARE name and need no sheet pin. 2. A child net merged with the root through a sheet pin takes the **root-side label**: `Project.add_sheet(child, ..., nets=["X"])` wires each sheet pin outward to a root local label X, so the net becomes /X. 3. A child-internal label becomes /<sheet>/NAME.

1.1 Global rails - power symbols, bare names

Net	Present on	Driven by	PWR.FLAG?
'GND'	all five	U1 RTN pin (passive)	**yes**
'V48_RAW'	poe, pwr	J1 V+ pin (passive)	**yes**
'V48_RTN'	poe	**only** J1 V- pin (passive)	**yes**
'+48V_SW'	pwr, expansion	U22 OUT	yes (eFuse output is a passive/power-out pin depending on the symbol)
'+12V'	pwr, expansion	forms after L20 (passive)	**yes**
'+3V3'	pwr, eth, mcu, expansion	forms after L21 (passive)	**yes**

All six PWR_FLAGS live on the sheet that owns the source (poe for V48_RAW/V48_RTN/GND, pwr for the rest), at #FLG0100+ / #FLG0300+. Power symbols being global, one flag each drives the net hierarchy-wide.

'V48_RTN' is deliberately confined to the **'poe'** sheet. It is the raw PoE negative, upstream of the hot-swap FET, and sits up to 57 V *below* board GND. Nothing outside the PD front end may touch it. If it ever appears on another sheet, that is a design error, not a routing convenience.

'V48_RAW' replaces both **'VPOE'** and **'+48V'** from **'research/power.json'**. The PD interface switches the *return*, not the positive rail, so there is only one positive 48 V node upstream of the eFuse. V48_SW from `research/interface-poe-48v.json` is renamed **'+48V_SW'** for consistency with the other rails.

1.2 Root-crossed signals - /NAME These are the names `constraints.json` uses, and they are contractual.

Final net	Exposed by	Members
**'/ETH_TXP'	poe + eth	J1 TD+, U10 TXP (pin 2), D10
**'/ETH_TXN'	poe + eth	J1 TD-, U10 TXN (pin 1), D10
**'/ETH_RXP'	poe + eth	J1 RD+, U10 RXP (pin 6), D10
**'/ETH_RXN'	poe + eth	J1 RD-, U10 RXN (pin 5), D10
 '/ETH_LED_LINK'	poe + eth	J1 LED cathode, U10 LINKLED (active low) via R7
 '/ETH_LED_ACT'	poe + eth	J1 LED cathode, U10 ACTLED (active low) via R8
 '/T2P'	poe + mcu	U1 T2P through its level-shift network, U30 GPIO47
**'/ETH_SCLK'	eth + mcu	U10 SCLK (pin 29), R31, U30 GPIO12
 '/ETH_MOSI'	eth + mcu	U10 MOSI (pin 28), R32, U30 GPIO11
 '/ETH_MISO'	eth + mcu	U10 MISO (pin 27), U30 GPIO13
 '/ETH_CS_n'	eth + mcu	U10 CS_n (pin 32), R33 pull-up, U30 GPIO10
 '/ETH_INT_n'	eth + mcu	U10 INT_n (pin 33), U30 GPIO14
 '/ETH_RST_n'	eth + mcu	U10 RST_n (pin 37), R34 pull-up, U30 GPIO21
 '/ENABLE'	pwr + mcu + expansion	U22 SHDN, R69 10k pull-down, J4 pin 23, U30 GPIO36
 '/FAULT'	pwr + mcu + expansion	U22 FLT (open drain), J4 pin 24, R132 10k pull-up, U30 GPIO37
 '/IMON'	pwr + mcu	U22 IMON, R68, U30 GPIO8
 '/PWM0' .. '/PWM7'	mcu + expansion	U30 GPIO4/5/6/7/15/16/35/38, J4 pins 1/2/5/6/7/8/11/12
 '/DSPI_SCK' ' /DSPI_MOSI' ' /DSPI_MISO' ' /DSPI_CS_n'	mcu + expansion	U30 GPIO39/40/41/42, J4 pins 14/15/16/17
 '/I2C_SCL' ' /I2C_SDA'	mcu + expansion	U30 GPIO17/18, R130/R131 4k7 pull-ups, J4 pins 18/19
 '/ADC0' ' /ADC1'	mcu + expansion	U30 GPIO1/2 via R136/R137 + D41/D42 clamps, J4 pins 20/21
 '/ID_ADC'	mcu + expansion	U30 GPIO9 via R135 + D40 clamp, R134 divider top, J4 pin 22

The five names in **bold** appear in `constraints.json` (`high-speed`, `diff_pairs`). Their sheet routing is therefore contractual: ****any change that stops one of them crossing the root renames it and breaks the constraint file****, and `check_return_path` raises `CheckError -> exit 2` on a `high-speed` net that is not on the board.

1.3 Sheet-internal signals - /<sheet>/NAME

Sheet	Final net	Members
poe	 '/poe/DEN_TAP'	R1/R2 junction, brought to a test pad - grounding it disables the PD
poe	 '/poe/CLS'	U1 CLS pin, R3 (the D-01 lever)
poe	 '/poe/CDB'	U1 CDB (converter disable, RTN-referenced) -> U20 EN
poe	 '/poe/SHIELD'	J1 shield tabs, R6 1M, C3 1nF/2kV hybrid to GND
eth	 '/eth/XI' ' /eth/X0'	Y10, C30, C31, U10 pins 30/31
eth	 '/eth/EXRES'	U10 pin 10, R30 12.4k 1 %
eth	 '/eth/1V20' ' /eth/TOCAP'	U10 pins 22/20 with C33 / C32
pwr	 '/pwr/SW'	U20 SW, L20 pin 1, D20 cathode, C54 low side
pwr	 '/pwr/BST'	U20 BST, C54 high side
pwr	 '/pwr/FB48'	R60/R61 junction -> U20 FB
pwr	 '/pwr/SW33'	U21 SW, L21 pin 1
pwr	 '/pwr/FB33'	R63/R64 junction -> U21 FB
pwr	 '/pwr/PGOOD'	U22 PGOOD -> D21 power-good LED via R71 (**no GPIO** - see mcu s4)
mcu	 '/mcu/EN'	U30 EN pad, R100 10k to +3V3, C85 1uF, J2 pin 5

```
| mcu | '/mcu/BOOT' | U30 GPIO0, R101 10k to +3V3, SW1, J2 pin 6 |
| mcu | '/mcu/TXD0', '/mcu/RXD0' | U30 GPIO43/44, J2 pins 3/4 |
| mcu | '/mcu/STATUS' | U30 GPIO48 -> R102 -> D30 |
```

2. Refdes allocation - unique ACROSS sheets, contractual for P4

constraints.json references U20 (thermal), J1/U30/J3/J4/H5 (edges), U20/L20/D20/U21/L21/U10/Y10/J1 (separation), and five group anchors. ****Renumbering any of those without editing constraints.json silently drops the constraint**** - place_anneal skips separation refs that are not on the board, and it does so without an error.

```
| Sheet | U | J | H | R | C | D | L | Y | SW | pwr_base |
|----|---|---|---|---|---|---|---|---|---|
| 'poe' | U1-U9 | J1 | R1-R19 | C1-C19 | D1-D9 | - | - | - | 100 |
| 'eth' | U10-U19 | - | R30-R59 | C30-C49 | D10-D19 | - | Y10 | - | 200 |
| 'pwr' | U20-U29 | - | R60-R99 | C50-C79 | D20-D29 | L20-L29 | - | - | 300 |
| 'mcu' | U30-U39 | J2 | R100-R129 | C80-C119 | D30-D39 | - | - | SW1 | 400 |
| 'expansion' | U40-U49 | J3, J4, H5 | R130-R159 | C120-C139 | D40-D49 | - | - | - |
500 |
```

2.1 poe sheet

```
| Ref | Part class | Note |
|----|---|---|
| J1 | RJ45 PoE magjack, THT right-angle shielded, integrated magnetics **and integrated bridge** | HanRun HY931147C class. **P3 must confirm from the datasheet that BOTH the data-pair centre taps and the spare pairs feed the internal rectifier** - a Mode-A-only or Mode-B-only jack is not 802.3 compliant. See 'decisions.md' OPEN-B |
| U1 | 802.3at Type 2 PD interface, SO-8-EP | TPS2378 class. **Pin 8 stays unconnected** so the TPS2379 second source drops into the same footprint |
| D1 | 58 V unidirectional TVS, SMB, 600 W | across 'V48_RAW'/'V48_RTN', physically first |
| C1 | 0.1 uF / 100 V X7R | VDD-VSS bypass; the standard's 50-120 nF window |
| C2, C4 | 22 uF / 100 V ceramic (x2 = 44 uF CBULK) | >= 5 uF for AC MPS, << 180 uF port ceiling |
| R1, R2 | 12.4 k 1 %, **0805 or larger** | split RDEN = 24.9 k with the tap out at '/poe/DEN_TAP' |
| R3 | **90.9 ohm 1 % (af) -> 63.4 ohm 1 % (at)**, 0603 | **THE D-01 LEVER.** Standalone, clearly silkscreened pad pair, no neighbours in the same silk box |
| R4, R5 | T2P level-shift network | T2P is RTN-referenced and reaches 57 V when high-Z - never a bare GPIO connection |
| R6, C3 | 1 Mohm + 1 nF / 2 kV | shield hybrid, with a fitted-by-default 0 ohm alternate footprint |
| R7, R8 | magjack LED series resistors | W5500 LED outputs are active-low, IOL >= 8.6 mA. **Confirm the jack's LED polarity at P3** |
```

2.2 eth sheet

```
| Ref | Part class | Note |
|----|---|---|
| U10 | W5500, LQFP-48 | RSVD pin 23 to GND; pins 38-42 NC; VBG (18) NC; DNC (7) NC; PMODE[2:0] NC |
| Y10 | 25 MHz, **CL 18 pF**, AT-cut fundamental, SMD3225-4P | select on the **total** ppm budget (initial + temp + ageing) staying under +/-50 ppm, not on the +/-30 ppm initial alone |
```

C30, C31	****27 pF COG****	'C = 2 x (CL - Cstray)', Cstray ~4 pF. ****Expect to trim on the first prototype.**** Most W5500 module schematics use 22 pF, which back-solves to CL 15 pF - wrong for an 18 pF part
R30	****12.4 k 1 %****	EXRES1 - sets the MDI drive current and therefore the 950-1050 mV output amplitude
C32	4.7 uF	TOCAP
C33	10 nF	1V20
C34-C40	100 nF x n + 4.7 uF bulk	VDD/AVDD decoupling
D10	4-channel low-capacitance TVS array, ****<= 1 pF/line****	PHY side of the magnetics, at the J1 end of the pairs. ****Fitted, not DNP****
R31, R32	22-33 ohm series, at the driver pin	'/ETH_SCLK' and '/ETH_MOSI', placed at U30's end. Footprint fitted so the value can be tuned if EMC bites
R33	10 k to +3V3	'/ETH_CS_n' pull-up. ****Mandatory****: GPIO10 has a 60 us low glitch at power-up and low = W5500 selected
R34	10 k to +3V3	'/ETH_RST_n' pull-up (belt and braces - the W5500's own is internal)

2.3 pwr sheet

Ref	Part class	Note
U20	100 V-class buck, 2 A, ESOP-8	SCT2A25 class. Exposed pad to In1 GND, >= 9 vias. 'EN' driven from U1's '/poe/CDB' for correct PoE start-up sequencing with no glue
L20	68 uH power inductor, Isat >= 3 A	~7 x 7 mm
D20	100 V Schottky, SMC	SS510 class. ****Dissipates more than U20**** - give the cathode >= 100 mm2 of F.Cu
C50-C53	2.2 uF / 100 V x2 in, 22 uF x2 out	100 V ceramics; ****no aluminium electrolytic on this board****
C54	BST cap	
R60, R61	FB divider for 12 V	
U21	3 A synchronous buck, 4.5-17 V, SOT-23-THIN-6	TPS563201 class. ****LDO disqualified****: 6.1 W at 0.7 A
L21	4.7 uH	
C55-C58, R63, R64	in/out caps + FB divider	
U22	60 V eFuse, HTSSOP-20-EP	TPS16630 class. ****MODE open = LATCH OFF.**** ILIM set to 1.0 A
****R69****	****10 k pull-down on '/ENABLE'****	****THE CAR-REQ-08 FAIL-SAFE.**** SHDN's open-circuit voltage is 2.48-3.3 V with a 10 uA source, so an unconnected SHDN floats HIGH and the device powers up ON. Must be called out on the schematic
R65-R68, C59	ILIM, UVLO/OVP divider (****0805 or larger****), IMON scaling, dV/dT	dV/dT sized ****fast**** - the daughter owns the inrush ramp
R70	****100 k bleed on '+48V_SW'****, 0805 or larger	de-energises the connector pins whenever ENABLE is low
D21, R71	green power-good LED from U22 PG00D	****hardware-driven, no GPIO****
D22, R72	red fault LED from U22 FLT	****hardware-driven, no GPIO****

2.4 mcu sheet

Ref	Part class	Note
U30	****ESP32-S3-WROOM-1-N8****	SKU frozen non-octal-PSRAM and non-R16V (GPIO35-37 and 47/48 are in use). Compatible: -N4, -N8R2, -N16R2
C80-C84	22 uF + 3x 100 nF + 1 uF	module supply must sustain a 500 mA Wi-Fi TX burst without browning out
J2	6-pin 2.54 mm header: GND, +3V3, TXD0, RXD0, EN, BOOT	Q9 default (b). ****Inside the enclosure, no panel cutout.**** Silkscreen: "PoE OFF or ISOLATED ADAPTER ONLY"
R100, C85	10 k + 1 uF EN RC	

R101, SW1	10 k pull-up + optional BOOT tactile	GPIO0
D30, R102	status LED from GPIO48	Q11 default: this plus the two magjack LEDs, all visible through the RJ45 cutout

Not on this sheet by design: no USB-C. A non-isolated PD achieves 802.3 compliance by having **only** the Ethernet connection and no accessible non-isolated conductor. ****Q9 option (a) "USB-C on every fixture" is therefore not available**** unless Q5 flips to isolated.

2.5 expansion sheet

Ref	Part class	Note
J3	2x7 2.54 mm THT male header, 250 V / 3 A	expansion **POWER** . Pin map: 'connector-icd.md' s3
J4	2x12 2.54 mm THT male header, 250 V / 3 A	expansion **SIGNAL** . Pin map: 'connector-icd.md' s3
H5	'MountingHole.3.2mm.M3'	CAR-REQ-15 support point between J3 and J4; also a keying feature
R130, R131	4.7 k I2C pull-ups	carrier-side, per requirements s2.1
R132	10 k '/FAULT' pull-up	open-drain net shared with U22's FLT
R134	ID divider top leg to +3V3	the daughter supplies the bottom leg (CAR-REQ-07 / Q10)
R135-R137, D40-D42	series resistors + 3.3 V clamps	on '/ID_ADC', '/ADC0', '/ADC1' **part of the CAR-REQ-14 survivability story, not just ESD hygiene** : a mis-seated daughter can bridge a neighbouring pin onto these

3. P4 generator notes (schlib specifics that bite)

- Six PWR_FLAGS.** Every rail on this board is driven by a *passive* pin - a connector pin, an inductor, an eFuse output. Without flags, ERC reports all six as undriven.
- 'V48_RTN' needs its own power symbol** and must appear only on poe. If schlib has no 48 V power symbol, use a generic power symbol with the name set explicitly - do **not** fall back to a local label, or the net becomes /poe/V48_RTN and constraints.json.voltages misses it (netlist_audit raises missing_net at error severity).
- Use the free-cluster 'hier_pin(net, at=...)' variant** for every net in sl.2 - each has 2+ components on it, and that variant places a local label plus the hierarchical label on one stub so the hier label joins by wire geometry rather than by name merging.
- Root-side labels sit 7.62 mm left of each sheet pin**, and sheet pins stack down the sheet symbol's left edge. Leave that strip clear and give sheets sharing a net the same nets= order.
- 'expect={...}' on U1, U10, U20, U21, U22, U30, J1, J3, J4** - pin-name insurance against a wrong symbol on the parts where a wrong pin is a dead board.
- Decoupling metadata:** Project.save(decoupling=...) must associate C80-C84 with U30's 3V3 pins, C34-C40 with U10's VDD/AVDD, C50/C51 with U20 VIN, C1 with U1 VDD (rail V48_RAW, gnd V48_RTN - **not** GND; this one needs an explicit gnd_net override).
- 'J2' silkscreen** must carry the non-isolated warning (s2.4).
- 'R3' silkscreen** must carry af=90R9 / at=63R4 so the D-01 lever is visible on the bare board.

4. Interface pinouts (silk-labelled, canonical)

- J1 magjack** (10 signal + 4 LED positions): TD+/TD-/RD+/RD- to /ETH_TXP /ETH_TXN /ETH_RXP /ETH_RXN; V+ (P9) -> V48_RAW; V- (P10) -> V48_RTN; shield tabs -> /poe/SHIELD. **P3 confirms P7/P8 from the datasheet** - they differ between HY931147C and HR861153C.
- J2 recovery header** (1x6, 2.54 mm): 1 GND, 2 +3V3, 3 TXD0, 4 RXD0, 5 EN, 6 BOOT.

- **J3 / J4 expansion:** `connector-icd.md` s3. That document is the ICD; this table must never disagree with it.

5. Board-edge placement (mirrors `constraints.json.placement.edges`)

Ref	Edge	pos	Why
‘J1’	top	0.21	the plug must seat through an enclosure cutout; 58 mm from the antenna, on a different edge
‘J2’	top	0.87	recovery header reachable with the lid off, far from the 48 V domain
‘U30’	right	0.50	antenna at a board edge, away from the shielded jack, and clear of both right-hand mounting holes
‘J3’	bottom	0.24	expansion power, near the eFuse and the DC-DC
‘H5’	bottom	0.46	CAR-REQ-15 support point between J3 and J4
‘J4’	bottom	0.72	expansion signal, near the module

Orientation is **not** specified via `rot`, because it depends on each footprint’s native orientation. Two P6 checks that must be done by looking, not by assuming (see `LEARNINGS.md` 2026-07-28 ”Connector mating direction”):

- **J1’s opening must face -y** (out of the top edge). Render an orthographic **side** view; the WRL bounding box is a known trap.
- **J3 and J4’s long axes must be parallel to the bottom edge**, pin 1 at the left end of each.

J3 and J4’s final board coordinates are an ICD deliverable. After P6, compare them against the nominal positions in `connector-icd.md` s7 and correct with `place_edit` if the annealer has moved them. **This must be done before any daughter run starts.**

Full architecture narrative (not inlined; see the artifact index): `architecture/blocks.md`, `architecture/decisions.md`, `architecture/power_tree.md`, `architecture/stackup.md`.

5 Schematic

Gate `erc` (phase P4): **pass** after 2 attempt(s); last 2026-07-28T20:56:57, failing 0 of 0.

Review waivers

P4 waivers and triage - LUMINA carrier (LUM-CAR-A)

Source: fresh-context `schematic-reviewer`, 19 findings against the ERC-passing schematic. ERC gate: **0 errors / 0 warnings** (first attempt). `netlist_audit`: **exit 0**.

Triage key: **FIXED** = changed in copper this phase | **WAIVED** = accepted, reasoning below | **AMENDED** = fixed by changing the ICD rather than the board | **OWNER** = raised at H2.

1. The most important caveat: what the ERC pass does and does not prove

Finding 7, WAIVED with a narrative correction that must survive into the design document.

`gen/lib_pin_types.py` retypes every pin in `lib/aiee.kicad.sym` to either `power_in` (where the datasheet extract says `ground/power_in`) or `passive` (everything else). Confirmed against the exported netlist: across all 102 components the only two pin electrical types present are `passive` and `power_in`. There is no `output`, `input`, `bidirectional`, `tri_state` or `open_collector` anywhere.

Consequence: KiCad's output-vs-output conflict, input-not-driven, and driver-conflict checks are **structurally inert** on this design. The 0-error/0-warning result proves only that

- no `power_in` pin lacks a driver, and
- no wire or label dangles.

It does **not** prove pin-level electrical correctness. That coverage came from the reviewer's independent per-IC audit against the datasheet extracts, not from ERC.

Why it is waived rather than fixed: the retype is the documented workaround for easyeda2kicad symbols carrying `unspecified` on nearly every pin, which otherwise floods ERC with false `pin_to_pin` warnings and false `pin_not_driven` errors and makes the gate unpassable (LEARN-INGS 2026-07-27 [easyeda2kicad][erc]). A richer retype is possible but would need per-pin direction data the extracts do not carry for every part.

Required of the design document: do not present ERC 0/0 as pin-conflict coverage.

2. Fixed in copper this phase

#	Severity	Item	Fix
2	warning	U22 '+48V_SW' had no negative-transient Schottky, against TPS16630 sec 9.4.1, while driving an off-board ~2800 uF inductive load. V(OUT) abs max is -0.3 V	SS510 added, cathode to '+48V_SW', anode to 'GND', adjacent to U22 pins 18-20
8	warning	R71 = 1 k did double duty as PGOOD pull-up (vendor range 10 k-100 k) and D21 ballast, giving 0.20-0.70 mA of LED current depending only on the green bin's 2.6-3.1 V Vf	PGOOD pull-up moved to 10 k; LED given its own ballast
10	note	'ENABLE' / 'FAULT' ran from the user-mateable J4 straight to the MCU and to U22 SHDN (abs max 5.5 V) with no series protection, while all three analogue connector signals had 1 k + clamp	1 k series added on both, mirroring the analogue pattern, without defeating R69
11	note	'/IMON' can exceed the ESP32-S3's 3.6 V pin abs max for ~1 us between fast-trip and the 45 A SCP threshold	1 k series added to IO8
12	note	W5500 AVDD pin 21 had only the shared 4.7 uF bulk while every sibling supply pin had a local 100 nF	100 nF added at pin 21
17	note	IO3 left floating, but the module datasheet states its strapping value "must be controlled by the external circuit that cannot be in a high impedance state" and it has no internal pull	10 k pull-down added
13	note	'TPD4E1U06DBVR' symbol labels pin 1 "D1+" / pin 4 "D2-"; the datasheet says pin 1 = D2-, pin 4 = D1+. Netlist was already electrically correct (wired by number) but the PDF would mislead layout and DFM	Symbol pin names corrected; recorded in 'lib/EDITS.md'

3. Fixed by amending the ICD instead of the board

#	Item	Amendment
6	'/FAULT'	also carries the carrier's own indicator LED, so a daughter asserting it sinks ~4.3 mA, not the ~0.33 mA the ICD implied - 13x what two other boards were designing against
3	The ICD described the average-energy governor as closed-loop on IMON, but IMON's specified accuracy starts at 0.6 A	and the rail's sustained limits

are ****0.25 A (af) / 0.50 A (at)**** - both below the floor | ****ICD rev A6**** adds s6.2.1: IMON is monotonic and useful as a guard, but is ****not**** a calibrated meter below 0.6 A. Daughters must not depend on better than $\sim \pm 20\%$ there without per-unit characterisation (the ID EEPROM already provides storage), a shunt+amp on a future revision, or a conservative governor |

4. Waived - accepted, with reasoning

Finding 5, U22 OVP set point sits above the part's own rating. OVP rising is 64.20 V typ and up to 66.78 V worst-case-high, against V(IN) recommended max 60 V and abs max 67 V. Between 60 V and ~ 66.8 V the eFuse runs outside its recommended range with neither OVP nor the TVS acting (SMBJ58A V_BR min is 64.4 V). **Waived because it is inherent, not a design error:** the OVP protects the **daughter**, not the eFuse. The TVS is the only real clamp on U22 itself. Recorded here and in the design document so nobody later reads "OVP" as protecting U22.

Finding 9, D30 status LED current is bin-dependent (2.12 mA down to 0.61 mA across the green Vf 2.6-3.1 V spread from a 3.3 V GPIO through 330R). Works; the current is set by the LED bin rather than the resistor. A red or a 2.0 V-class green would be predictable. Not worth a respin.

Finding 14, J1 pin 3 (RX centre tap) floating. Neither the W5500 nor the HY931147C datasheet publishes a centre-tap network, so this could not be grounded from either extract; the W5500 biases RXP/RXN internally. Raised at H2 as a cheap optional lever (a 0.1 uF DNF footprint to GND) since P5 freezes the layout.

Finding 15, no crystal drive-level damping. R36 = 0R gives the oscillator no series damping against the W5500's specified 59.12 uW drive level. R36 is the intended trim lever and is fitted; measure on the first prototype. Crystal loading is otherwise correct (27 pF back-solves to ~ 4.5 pF stray for an 18 pF part).

Finding 18, MPS depends on the W5500 staying awake. Resistive + bias load alone is ~ 4.6 mA; normal operation adds ~ 5.3 -10.8 mA (W5500) and ~ 3.3 -8.2 mA (ESP32), so MPS is comfortably over the 10 mA DC floor. But PHY power-down **plus** ESP32 deep sleep would fall to ~ 6 mA and the PSE would drop the port. Already captured as firmware requirement **FW-05** in ICD s7.8.

Finding 19, documentation drift. (a) ICD header rev label - fixed, now carries the full A2-A6 history. (b) `power_tree.md` says CBULK 44 uF (2 x 22 uF); the build fits 4 x 10 uF = 40 uF nominal because no 22 uF/100 V MLCC exists on LCSC. Electrically fine (≥ 5 uF AC-MPS floor, far under the ~ 180 uF port ceiling); the doc will be reconciled in the design document.

netlist_audit 'diffpair_naming' x3. /ETH_TXP+/ETH_TXN and /ETH_RXP+/ETH_RXN do not use the _P/_N suffix convention the auto-discovery heuristic looks for. Cosmetic only: `constraints.json` declares `diff_pairs` **explicitly**, so the heuristic is never consulted and the 100 ohm targets bind correctly.

5. Raised at H2 for the owner

Findings 4 + 5 are coupled and cannot both be comfortably satisfied with this part. Re-derived independently from the TPS16630's asymmetric limits with 1 % resistors:

	Threshold	worst-low	typ	worst-high	Requirement	Result
	UVLO rising	33.66 V	35.02 V	**36.42 V**	< 37 V	**PASS** , 1.6 % margin
	OVP rising	61.69 V	64.20 V	**66.78 V**	< 67 V abs max	pass, but above the 60 V rec. max
	OVP falling	**57.18 V**	60.03 V	-	> 57 V	**PASS** , but only **0.18 V**

The 0.18 V OVP-falling margin does not survive thick-film tempco: at +/-100 ppm/degC over a 60 degC excursion the uncorrelated worst case is ~1.2 % on the ratio, i.e. ~56.5 V - below the 57 V PSE maximum. The consequence is bounded (the rail fails to re-enable only after a genuine >61.7 V event while the PSE then sits at exactly 57 V), but it is not a margin.

The constraint is close to unsatisfiable with this part: the OVP hysteresis plus +/-2 % forces $V_{OV_fall_min} \leq V_{OV_rise_min} / 1.078$, so "never false-trip below 57 V" and "guaranteed re-enable at 57 V" cannot both be held with comfort. **Owner choice: (a) accept and document the corner, (b) specify 0.1 % resistors on R66/R67/R73, or (c) re-centre toward protecting U22 and accept that the rail may not re-enable at exactly 57 V. **Recommendation:** (b)** - it is a BOM tolerance change, not a respin, and it buys the margin outright.

Finding 16, the class-4 upgrade resistor has zero thermal margin. The fitted R3 = 90.9R (class 3, af) dissipates 73 mW on a 0603 - fine. The documented class-4 alternate (63.4R) dissipates **105 mW against a 100 mW 0603 rating**. Classification is a sub-second transient so it will survive, but D-01 promises the at upgrade is "a resistor change only - no respin". **Specify the alternate as 0805** to keep that promise clean. Related: with class 3 programmed, a Type-2 PSE never runs 2-event classification, so /T2P reads "no Type-2" permanently on build 1 - expected, and firmware must not treat it as a fault.

Finding 1 (the only ERROR) is being fixed, not waived: parts.json had drifted badly from the schematic - 9 refdes with no BOM line, 2 with the wrong part, 2 phantom lines (the U40 EEPROM superseded by the ID_ADC scheme, and a phantom C120), and 7 wrong quantities. Being resynced from the exported netlist so the BOM/CPL that P9/P10 emit is correct.

reports/schematic.pdf not available.

6 Layout

Gate place (phase P6): **pass** after 1 attempt(s); last 2026-07-29T11:02:27, failing 0 of 0.

Gate drc_routed (phase P7): **pass** after 2 attempt(s); last 2026-07-29T18:55:29, failing 0 of 0.
No board renders found.

7 Verification

Gate verify (phase P8): **fail** after 2 attempt(s); last 2026-07-29T22:11:27, failing 114 of 118.
reports/verify_all.json not available.

8 DFM and Fabrication

Pending — produced at P9.

9 Run Record (Appendix)

Phase digests

P0

P0 Intake digest - LUMINA carrier (LUM-CAR-A)

- Inputs: 3 briefs ('00' system context v2.0, '01' carrier brief, '05' closed decisions).
- '05' is binding and supersedes '00' section 6 for D-01/D-02/D-03.
- Artifact: 'architecture/requirements.md' (requirements-analyst, spawned as subagent).
- Function: universal PoE carrier. One Ethernet cable in; fixture-agnostic expansion

interface out. PD front end + 48->12->3.3 V + ESP32-S3 + W5500 + RJ45 magjack.
 No LED drivers, no light-engine storage, no audio, no wireless control path.
 - Closed on entry: D-01 (at-sized power stage, af classification, resistor-only upgrade),
 D-02 (connector carries 48 V raw + 12 V + 3.3 V), D-03 (no UV board).
 D-04 (strobe colour) is NOT this board's - carrier exposes the full 8-PWM ceiling so
 either answer works.
 - Hard ceiling found in the brief and carried forward: ESP32-S3 LEDC = 8 channels,
 4 timers, low-speed only, 14-bit max. Default 13-bit @ 9.77 kHz. The legacy "16-bit @ 1.2 kHz" is not achievable and is recorded as dead.
 - 16 open questions raised. Two are SAFETY-flagged and will not be guessed past H1:
 Q5 converter isolation + enclosure material, Q6 worst-case per-rail connector current.
 - Orchestrator is a background agent with no human channel. Decision
 D-PO-PROVISIONAL:
 proceed P1/P2 on the analyst's stated defaults, marked PROVISIONAL throughout;
 H1 is the single blocking confirmation point. If the human overrides Q5 or Q6,
 P2 architecture must be revised before P3.
 - Blocking for P5 (not yet reached): MECH-02 board outline in mm is permanent at
 'board_init.py --outline WxH'. Must be closed with the human at H1.

P1

P1 Research digest - LUMINA carrier (LUM-CAR-A)

Roster: 8 parallel agents - power-architect, refdesign-poe-pd,
 interface-poe-48v,
 interface-ethernet, and 4 component scouts (connector-expansion, poe-power,
 mcu-net, rail-protect). All 8 fragments written and JSON-validated.

Findings that change the brief

- **The brief's named PD part is disqualified.** Si3402-B/-C AND Si3404 are IEEE
 802.3 Type 1 only (no Class 4 at any resistor value); Type 2 additionally needs
 2-event classification *recognition*, which is silicon. D-01's "resistor
 change,
 no respin" is unachievable with them. Skyworks' actual PoE+ parts are
 effectively
 unstocked on JLC. Consequence: the "~10 W regulated" figure in requirements
 3.2
 comes from AN956 and describes a part this board cannot use - re-derive, do not
 inherit.
 - **Lead pair: TPS2378DDAR (PD interface) + SCT2A25STER (100 V buck), ~\$2.70
 @30.**
 Class program 90.9R (class 3, af) -> 63.4R (class 4, at) = one 0603 change.
 Every at-capable PD IC on JLC that integrates a converter is transformer-based,
 so a non-isolated buck at Type 2 REQUIRES the two-chip solution.
 - **48 V raw at "2 A continuous" is not deliverable by any PSE.** 2 A at 48 V
 is
 96 W on a 25.5 W supply. Resolution proposed: 2-3 A stays the connector PIN
 rating; load switch limits at 1.0 A with latch-off; ICD publishes 0.25 A (af) /
 0.5 A (at) sustained.
 - **The 48 V load switch is a COMPLIANCE requirement, not just protection.**
 802.3 caps PD port capacitance near 180 uF; the strobe daughter holds ~2800 uF,

so the switch must be off through PD power-up.

- ****Carrier overhead is higher than the brief's 1.5 W:** 2.44 W (af) / 3.75 W (at)**
computed; the brief's chain double-counts regulator loss. ESP32-S3 + W5500 alone measure 0.70-0.76 W before converter losses.
- ****HR911105A and all 7 other high-stock HanRun jacks are NOT PoE-capable.****
Only 4 PoE magjacks exist on LCSC. All three HanRun PoE jacks are JLC-assemblable (wave soldering), so THT is not an assembly blocker.
- ****Fine-pitch mezzanine is closed by CAR-REQ-17.**** Every 0.4-1.0 mm family JLC stocks rates 50-60 V. Pitch floor is ≥ 2.00 mm. 2.54 mm THT is the viable class.
- ****The 15 mm standoff default is unachievable**** (stocked 2.54 mm parts mate at 11.0 mm), and at 11 mm a board-edge THT RJ45 (~13-16 mm tall) collides with the daughter board.

Binding numbers for later phases

- Stackup: ****4-layer JLC04161H-3313 required**** - the only stackup with a diff_100 profile (0.260/0.210). 2-layer computes to 1.081 mm per leg. Also independently forced by thermal (buck 1.35 W at dt_c 70 passes on 4L, fails on 2L).
- 48 V spacing: ****0.60 mm outer / 0.10 mm inner**** (IPC-2221B B2/B1, 51-100 V band). 57 V < the IEC 62368-1 ES1 limit of 60 V, so this is functional insulation only.
- MDI: 2 pairs, 100 ohm diff, t_rise 3.0 ns, GND reference, max_skew 2.5 mm.
- W5500 SPI ceiling is ****33.3 MHz guaranteed**** (80 MHz is "theoretical design speed"). Recommend 20 MHz on SPI2 IO_MUX pins.
- ESP32-S3 forbidden pins: GPIO0/3/45/46 (strapping), GPIO19/20 (USB), GPIO35-37 on octal-PSRAM SKUs. Legal set proposed: SCLK 12 / MOSI 11 / MISO 13 / CS 10 / INT 14 / RST 21 with a 10k CS pull-up.
- 48 V and logic ****share one ground****. The magjack's 1500 Vrms protects the PHY from the power, not the board from the cable - PoE taps the cable-side centre taps, bypassing the barrier by design.

Two verified pipeline traps (recorded as decisions)

1. 'rules_gen.py' never reads the 'voltages' key - the 0.60 mm HV clearance is not enforced during routing and only surfaces at P8. Needs an explicit HV net class or '.kicad.dru' rule at P5.
2. 'planes_gen' has no keepout key, so the magjack plane void must come from plane 'region' rectangles, or P8 'check_return_path' fails unwaivably.

Orchestrator process note

One decision record ("Power topology") was written before the corresponding scout reported and without reading its file. The part numbers happened to match, but the efficiency figures in it were not verified. A CORRECTION entry supersedes it; the verified basis is 'research/poe-power.md' section 0.

P2

P2 Architecture digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'blocks.md', 'power_tree.md', 'stackup.md', 'sheets.md',
'constraints.json', 'connector-icd.md' (ICD-01), 'decisions.md'.

- ****Chain:**** PoE magjack (integrated bridge) -> TPS2378-class PD interface -> 100 V buck to 12 V -> sync buck to 3.3 V; 60 V eFuse gates 48 V to the daughter;
W5500 + 25 MHz crystal on SPI2 IO_MUX at 20 MHz; ESP32-S3-WROOM-1-N8;
split 2x7 power + 2x12 signal expansion connectors.
- ****Stackup:**** JLC04161H-3313, 4 layer, 1 oz. Forced by 100 ohm MDI (2L needs 1.081 mm/leg), not by thermals.
- ****Outline (MECH-02 proposal):**** 100 x 80 mm, 3 mm radius, 4x M3 at 5 mm inset plus a 5th M3 at (46,74) for CAR-REQ-15. From 4410 mm² of block area at 55-60 % utilisation; 80x60 needs 92 % and is impossible.
- ****Budget re-derived**** from the chosen parts, because the brief's "10 W regulated" describes a disqualified part: ****8.6-9.3 W (af) / 18.7-20.0 W (at)**** to the daughter. D-01's 8.5/18.5 W holds with margin. Carrier overhead 2.4/3.7 W.
- ****Cost:**** ~\$26-32 per assembled carrier at qty 14 against a \$30 provisional target. 12 carriers is ~\$360 of a \$500-1000 programme budget that also owes daughters, enclosures and a PoE switch.
- ****Four sign-off gates drafted:**** two-column af/at budget; PWM 4+4 across LEDC timers 0/1 (works for either D-04 answer); ENABLE fail-safe (one net, one 10k pull-down, on the only GPIO band with no power-up glitch); full 28-pin ESP32-S3 legality map.
- ****constraints.json validated:**** schema-clean (5 high_speed, 5 power, 2 diff_pairs, 6 voltages, 1 thermal, 2 planes, placement edges/groups/separation).

Deliberate omissions with a mandatory P5 recipe

'placement.keepouts' and 'planes[].region' are absent ****on purpose****:
'board_init'
does not origin the outline at (0,0) - the origin comes from the packed component
bbox - so absolute rectangles cannot be authored at P2. 'stackup.md' s7.1 carries
the exact P5 recipe: after 'board_init', read
'reports/board_init.json.outline.bbox'
and patch 'kicad/constraints.json' with the antenna keepout and the six plane regions. ****Skipping it leaves the ESP32-S3 antenna over a solid GND plane.****

Three pipeline traps carried forward

1. 'rules.gen' never reads 'voltages' (grep-verified) - nothing makes the P7 router honour 0.60 mm 48 V clearance. Add a named '.kicad.dru' rule keyed on 'A.NetName' ('A.Net' silently matches nothing) at P5.
2. 'planes.gen' has no keepout key - the magjack plane void must be built from plane 'region' rectangles or P8 'check_return_path' fails unwaivably.
3. 'rules.gen' puts every power net in ONE Power netclass at the widest width, so
+12V's 1.10 mm would be applied to V48-RTN's 0.6 A stub and to Freerouting's DSN - split 'netclass.patterns' before 'route_auto'.

Riskiest decisions

1. ****No LCSC PoE magjack publishes an 802.3at tap rating.**** HR871150C's own 17.5 W two-tap figure disqualifies it; the stock-safe integrated-bridge part is only af-rated. D-01's at upgrade stays a board-level resistor change but acquires two non-board dependencies: magjack qualification and enclosure vents.
2. ****The outline is permanent and two unanswered questions move it**** - Q8 (radio) and Q4/Q12 (stack height vs the ~15 mm-tall RJ45). Recommendation: 11 mm standoff plus a 30x26 mm top-edge notch in every daughter, which doubles as a mechanical interlock against 180-degree mating.

P3

P3 Parts + Library digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'parts/', 'lib/aiee.kicad_sym', 'lib/aiee.pretty', 'reports/lib_pull.json', 'reports/fp_verify*.json' + 'fp*.overlay.svg' (9 footprints verified against datasheet pad geometry), 'work/ps*.json', 'work/pull*.json'.

- ****PD controller: TPS2378-class, NOT Si3402-B.**** The briefs name Skyworks Si3402-B/AN956; both Si3402-B and Si3404 are IEEE 802.3 ****Type 1 only**** at any resistor value, so D-01's "resistor change, no respin" upgrade to Type 2 is unachievable with them. TPS2378 reaches Type 1/Type 2 on a ****single**** RCLS (90.9R af -> 63.4R at, 1 %). TPS2372/2373 rejected: two class resistors = a two-part upgrade. Consequence recorded: AN956's "~10 W regulated" figure is a Type-1-only statement and does not describe this design.
- ****Magjack: LINK-PP LPJG0926HENL (C22457393)****, after the first choice failed qualification. Screening history matters because it changed the answer twice:
 - HY/original 10/100 part: publishes ****no**** tap current rating at all.
 - Best documented 10/100 PoE+ part (Würth 7499410213) publishes 600 mA/tap = exactly the 802.3at DC maximum and ****below**** the 0.686 A peak; +7.20 USD/board, ****zero**** LCSC stock (DNP + hand-solder x14), internal Bob Smith network that cannot be removed, isolation pad gap collapsing 3.58 -> 1.05 mm.
 - Re-running the screen over ****gigabit**** ICMs (the earlier category rejection was an error) found LPJG0926HENL: ****720 mA @ 57 VDC continuous****, 1.20x the at DC max and ****1.05x the 0.686 A peak**** - the first candidate that covers the peak. 3109 in stock, 3.7852 USD at qty-10 (52.99 USD for 14, ****80.57 USD cheaper**** than the Würth plan), and JLC Assembly = Wave Soldering so JLC ****places**** it.
 - Accepted costs: no integrated bridge (two external bridges), return loss -16 vs -18 dB, CMRR -30 vs -35 dB, and ****0..+70 C**** op temp makes it the lowest-temperature-rated part on the board (+15.9 K margin at the ICD's 30 C room ceiling for at).
 - ****OVP divider on precision thin-film**** (R66/R67/R73 = Yageo RT 0.1 %, ****25 ppm/C****). Tempco, not initial tolerance, was the defect: 100 ppm/C over a 60 C excursion drags the OVP ****falling**** threshold to ~56.5 V, i.e. ****below**** the 57 V legal PSE maximum, so a legal rail could fail to re-enable after an overvoltage event. 25 ppm/C moves worst-case falling to ~58.2 V: ~1.2 V of real margin instead of 0.18 V. R4/R5 (T2P level shift) deliberately stay 1 %.
 - ****Three library defects fixed that no later phase could have resolved**** (all logged in 'lib/EDITS.md'):
 1. J1 board-lock <-> VC creepage was ****0.487 mm**** against the board's own 0.635 mm rule - voltage-derived, so P8 would have failed it after routing.

Pads 19/20 made oval 2.00 x 2.60 and VC1/VC4 1.20 mm: gap now **0.687/0.697 mm**
 with every annular ring at JLC's 0.15 mm minimum. Side effect: the chip-side/line-side barrier improved 1.401 -> 1.451 mm, clearing HALO's 1.40 mm guidance.
 2. R3's land moved 0603 -> **0805** so the 63.4R class-4 upgrade at ~105 mW is in-spec on the SAME land - D-01's no-respin promise made real.
 3. D1 SMBJ58A courtyard was self-intersecting from zero-length 'fp_line' segments (easyeda2kicad artifact); 5 removed across 2 footprints.
 - **BOM sync hardened:** 'bom_sync.py' now preserves 'not_fitted' lines. The first
 sync silently dropped C334927, the 63.4R class-4 resistor, precisely because it has zero refs by design - deleting the one component that makes the upgrade promise real. Final: 50 BOM lines, 109 placed components, 0 lines changed on a
 repeat run, 0 netlist LCSC codes without a BOM line. Only H5 lacks an LCSC property (mounting hole, board-only by design).

Process note of record

The orchestrator's own Task brief for the magjack swap carried the **Wurth** part's
 pin deltas into a brief for the **LINK-PP** - a different part with a different pin
 count. The agent detected the contradiction, built to the datasheet (VC1-4 = pins
 11-14, LEDs 15-18, EH 19-20, confirmed against 'parts/C22457393.json') and reported
 the override rather than splitting the difference. That behaviour is what kept a
 wrong magjack pinout from flipping the connector end-for-end.

P4

P4 Schematic digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'kicad/lumina-carrier.kicad_sch' + 'poe/pwr/eth/mcu/expansion' sheets,
 'reports/lumina-carrier-schematic.pdf', 'reports/erc-waivers.md',
 'work/erc_final.json', 'work/netlist_audit.json'.

- **Gate 'erc': PASS, 0 errors / 0 warnings** (2 attempts recorded).
 - **Netlist audit clean**; 109 placed components, 50 BOM lines.
 - **Reviewer (fresh context) returned 19 findings** and independently verified all
 ten known-critical items as correct: R69 ENABLE fail-safe, TPS2378 EP-to-VSS, APD-to-RTN, the 1.2 V divider, the EN abs-max chain, W5500 RSVD asymmetry, GPIO45, the magjack map, and both frozen ICD pin maps.
 - **Triage: 1 error to the fix loop, 7 fixed in copper, 2 fixed by ICD amendment
 (rev A6), 7 waived, 2 raised for the owner at H2.**

The most important finding is an HONESTY item, not a defect

Because 'lib_pin_types' types every non-supply pin 'passive', KiCad's output-vs-output, input-not-driven and driver-conflict checks are structurally
 inert on this schematic. The ERC 0/0 therefore proves only that no
 'power_in'
 pin lacks a driver and no wire dangles. **Real pin-level coverage came from the

reviewer's per-IC datasheet audit, not from ERC**, and the design document must not claim otherwise.

Two findings fixed by amending the ICD rather than the board

- **FAULT sink current published as ~4.3 mA** (the schematic implied 0.33 mA). Publishing the real number is cheaper for daughters than a carrier respin plus a GPIO.

- **IMON governor claim softened**: the eFuse current-monitor pin's specified accuracy starts at 0.6 A while the rail's sustained limits are 0.25/0.5 A - both *below* the accuracy floor, so the closed-loop governor claim needed qualifying rather than asserting.

H2 verdict: approved

Owner delegated the remaining decisions. Two items went to the owner and both were resolved in favour of the recommendation: the 0.1 % OVP divider (taken, and upgraded to thin-film for tempco - see P3) and the magjack choice (OPEN-A closed on LPJG0926HENL).

P5

P5 Board Setup digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'kicad/lumina-carrier.kicad.pcb', 'kicad/lumina-carrier.kicad.pro', 'kicad/lumina-carrier.kicad.dru', 'work/board.init*.json', 'kicad/constraints.json' (patched with real outline-derived rectangles), 'kicad/decoupling.json'.

- **Outline 100.0 x 80.0 mm**, read back from the report rather than assumed: 'outline_bbox' = [19.58, 57.132, 119.58, 137.132]. 'corner_radius' **3.0 mm honoured** (not clamped), 'mounting_holes' 4 x M3 plus the 5th at (46,74) for CAR-REQ-15. MECH-01 and the MECH-02 H1 proposal both satisfied.

- **Stackup JLC04161H-3313**, 4 layer, 1 oz outer. **In1.Cu = GND, In2.Cu = +3V3**,

each as three rectangles that route around the mounting holes; the ESP32-S3 antenna keepout is authored from the real bbox per the P2 recipe (skipping it would have left the antenna over solid GND).

- **Setup violations 92 -> 36**, and the 36 are all pre-placement artifacts (27 'copper_edge_clearance', 3 'courtyards_overlap', 3 mask bridges, 1 npth-in-courtyard from components still stacked at the origin) plus 2 benign warnings. **Zero library defects remained.**

- **Named '.kicad_dru' HV rules added** (P2 trap 1: 'rules_gen' never reads 'voltages', so nothing otherwise makes the router honour 48 V clearance): 'HV_48V_clearance', 'HV_48V_to_HV_48V', 'HV_48V_raw_to_rtn' at 0.635 mm, keyed on 'A.NetName' ('A.Net' silently matches nothing), plus 'magjack_isolation_barrier' at 1.30 mm.

- **Netclasses split** before routing (P2 trap 3) so +12V's 1.10 mm is not applied to V48_RTN's 0.6 A stub or pushed into Freerouting's DSN.

Three library defects found here that placement could never have resolved

1. **ESP32-S3 thermal-land vias-in-pad** at 0.075 mm ring / 0.25 mm drill, failing

both min-annular and min-hole. Enlarging them to 0.60/0.30 fixed those two and created ****24 clearance + 24 mask-bridge**** errors instead, because the via pads then touched the thermal sub-pads. They were ****removed entirely****: thermal vias under a ground land belong to the board, not the footprint - P7 stitching places them against the real GND pour with correct net assignment, and that also removes the via-in-pad solder-wicking risk already on the P9 list.

2. ****D1 SMBJ58A self-intersecting courtyard**** from zero-length 'fp_line' segments (easyeda2kicad artifact). Scanned board-wide: 5 removed across 2 footprints.

3. Sidecars ('constraints.json', 'decoupling.json', 'parts.json') relocated beside the board so every later script resolves them.

P6

P6 Placement digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'reports/gate-place.json', 'reports/place_seed.json', 'reports/place_anneal.json', 'reports/place_metrics_final.json', 'reports/place_region_audit.json', 'reports/place/*.ops.json', 'reports/render_place/*.png', 'reports/route-probe-p6.json'.

- ****Gate 'place': PASS, 0 violations**** (1 attempt).
- ****HPWL 4212 -> 3601 mm****, crossings ****955 -> 613****, utilisation ****24.7 %****.
- ****Route probe completion 0.9842**** (5 unrouted of 317), above the 0.98 bar - i.e. the placement was confirmed routable before P7 was entered.
- 116 footprints placed (109 schematic components + 5 board-only mounting holes, plus fiducial/board items).

Two self-inflicted issues found and fixed here

1. ****My own '.kicad.dru' HV rules were over-scoped.**** With no pad-pair exclusion they fired ****30 times between adjacent pins of the SAME package****: U1's own VDD/VSS at 0.200 mm, U22's HTSSOP pitch, and C1/C61/C62/C63's own two pads at 0.590 mm. No placement can fix a 0.65 mm-pitch package, and that was never the rule's purpose - TI's 0.635 mm guidance governs ****board copper around**** the part, not the vendor's lead frame. Added '!(A.Type=='Pad' && B.Type=='Pad')' plus a per-refdes same-courtyard exclusion (enumerated, never wildcarded) to all three rules: ****30 -> 0****. This loses no coverage because P8 'check_creepage' independently models pad-to-pad from the 'voltages' table.

2. ****A blanket silk-clip sweep was over-aggressive and was caught before commit.****

The first pass was scoped board-wide and removed ****241**** silk primitives, stripping legitimate silkscreen from nearly every passive - a 0603's body outline naturally overlaps its own pad bboxes. Caught by sanity-checking the count, restored from backup, and re-scoped to only the refs DRC actually flagged: ****exactly 5 clipped**** (J1 x4, D22 x1), verified by differencing the silk-layer references against the backup.

Debt this phase created, and it is real

Silk DRC violations were taken ****236 -> 5**** by 9 ops files carrying ****160**

'move_text' ops** that pushed refdes *outward* to clear overlaps. Nothing pulled them back. The board was therefore optimised for a zero-warning gate at the cost of assembly legibility: baseline **median refdes offset 4.079 mm, max 12.955 mm, and 100 of 116 refdes more than 1 mm beyond their own part's pad extent** (worst: C2 at +10.067 mm). The owner raised this as a requirement and it is repaid in P8, as a constrained minimisation - pull each label back toward its own part while 'silk_overlap' / 'silk_over_copper' / 'silk_edge_clearance' stay at zero.

P7

P7 Routing digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'work/p7/*' (DRC ladder, diff-pair and barrier measurements, plane repairs), 'work/p8/drc_confirm.json' (the standing 0/0 proof), 'kicad/lumina-carrier.kicad.pcb'.

- **Gate 'drc_routed': PASS, err+warn = 0** (1 attempt, ~3089 track segments, 4 layers, In1 = GND / In2 = +3V3 planes with stitching).
- MDI routed as two matched pairs; final 'check_diffpair' **0 violations**:

TX skew **0.0012 mm (0.008 ps)**, RX skew **0.98 mm (6.58 ps)** as the gate measures it (see the caveat below - the gate under-measures RX).

- Worst measured pad-to-track clearance on the HV domain **1.4221 mm**.

Measurement record settled (shape-aware), including a retraction

Three different numbers were produced for one barrier because a rectangle model was twice applied to circular pads. Definitive figures come from a single shape-aware tool ('work/p7/pad_gap.py'; capsule model: circle = degenerate segment, oval = stadium spine, rect = per-axis):

- **J1 chip-side/line-side barrier = 1.6130 mm** as built (1.4510 mm at pad 2's original size). Both pass the 1.30 mm DRU rule and HALO's 1.40 mm guidance. The original P4 figure of 1.451 mm was **correct all along**; the P7 "correction" to 1.148 mm was wrong and is retracted.
- **Shield board-lock to PoE tap = 0.6029 mm** (pad 19 to pad 11) and **0.6127 mm** (pad 20 to pad 14). These **pass IPC-2221B B2** (0.60 mm for the 51-100 V band) and fall 0.032 mm short only of the 0.635 mm adopted board-wide from TI's guidance.
- **Deliberate non-reversion:** J1 pad 2 stays at 1.200 mm although the shrink that produced it was based on the wrong model. 0.150 mm ring on a 0.900 mm drill is exactly JLC's PTH minimum, the barrier is *better* as built, and enlarging the pad now would shrink every pad-to-track gap around it on a board just brought to zero clearance errors. Recorded so the next reader knows this is an accepted state, not an unexamined leftover.

The expensive lesson of this phase

A 0.300 mm placement nudge of C35 - made to open a claimed 0.681 -> 1.681 mm GND fan-out corridor for U10 - ****introduced a short****. Two compounding errors: (1) the corridor figure read C35's ***footprint origin*** y, not its ***pad-1 centre*** y, off by exactly 0.700 mm, so the real corridor was 0.9811 mm; and (2) it was validated with 'gate place' (PASS) and 'check_decoupling' (unchanged) - ****both of which are structurally blind to routed copper**** - while at the new position C35's pad 1 landed on an existing /ETH_RSTn run, giving +6 DRC errors including a 'SHORTING_ITEMS' and a 'solder_mask_bridge'. The router caught it and reverted. The corridor could never have closed the item anyway: the real blocker was a ****void in the In1 GND fill**** created by neighbouring +3V3 plane vias' clearance holes, so a via there would have connected to nothing (0 legal 0.5 mm centres across the whole corridor at 0.025 mm step). The router solved it properly by escaping U10 pad 9 ****north**** into the LQFP's own die shadow.

****Rule of record:** on a routed board, DRC is the authoritative oracle for any placement or silk change.****** This is in 'LEARNINGS.md'.

Caveat carried into P8

'check_diffpair''s 'matched_terminals' pairs a p-pad with an n-pad only within 'TERM_PAIR_MM = 2.5'. J1's RX pads are ****4.58 mm**** apart, so ****J1 is not recognised as a terminal**** and the "trunk" is measured from the ESD array to the PHY only. The RX pair therefore ****passes at 0.98 mm**** while the true magjack-pad-to-PHY-pad skew is ****6.451 mm (43.0 ps)****, with ****4.326 mm** of it inside the J1 escape the gate never inspects******. This is an open owner decision at H4, not a waiver.

P8

P8 Verification digest - LUMINA carrier (LUM-CAR-A)

Artifacts: 'reports/checks/*.json' (verify_all, 8 checks),
'reports/render_silk/*.png',
'work/p8/creepage_adjudication.json', 'work/p8/ipc_adjudicate.json',
'work/p8/silk/refdes_{before,after}.json', 'work/p8/{tapcreep,hvpwr,silk}/',
'reports/review-board.{md,json}'.

- ****'drc_routed' re-confirmed 0 errors / 0 warnings**** after every single copper and silk edit in this phase, and again at H4. It never left 0/0.

- ****'verify' gate: 118 findings**** (114 error-severity + 4 warnings).
'check_current' 94, 'check_creepage' 11, 'check_return_path' 11 (9 error + 2 warning),
'check_decoupling' 2 (both warning). ****'check_diffpair', 'check.thermal', 'check_silk' and 'check_pdn' are all 0.****

Four real defects fixed, none of which any gate reported

1. ****A-side PoE tap differential creepage, 0.3292 -> 0.6502 mm**** (F.Cu; 0.6524 mm on the other three layers). '/poe/POE_TAP_A1' and '/poe/POE_TAP_A2' are the two AC inputs of

bridge D2, so the full line potential sits between them - but both are declared 57 V,
'check_creepage' computes 'dv = 0', and the pair is **never evaluated**. Found by hand.

Fixed by moving both transition vias (0.6 -> 0.5 mm, annular 0.1 mm = the board and JLC

minimum) to a grid-searched optimum that maximises the **worst** margin rather than the

headline gap - the first optimum bought a 0.726 mm tap gap while leaving +0.005 mm to

SHIELD, i.e. robbed one 57 V gap to pay another.

2. **A SECOND** tap pair at 0.5500 mm, on the same two nets, hidden behind the first and only

visible after it was fixed. Also fixed.

3. **Four J1 SHIELD-to-tap** creepage errors, previously classified as package-internal but

actually routed copper: 0.213 -> 0.6390, 0.2142 -> 0.6388, 0.220 -> 1.7700, 0.2129 -> 0.6620 mm, by replacing a 39-segment SHIELD staircase with a 6-segment polyline

at the analytic optimum. The corridor has **0.0778 mm** of total slack, so the topology was

forced, not chosen. Re-swept: 0 of 31 SHIELD items now under 0.600 mm.

4. **Two 0.200 mm** power-trunk segments now 0.500 mm (V48_RAW, 1.0 A). All five candidates

were first proven to be **bridges** in their net's connectivity graph, so each really does

carry the whole rail - no parallelism was assumed.

A new '.kicad.dru' rule, **'poe_tap_differential_pair'** (0.60 mm), now holds defect 1 from

recurring, because 'check_creepage' structurally cannot. It was **proven live** rather than

assumed: raised to 0.90 mm it fires 17 times and independently reports 'actual 0.6502' /

'0.6524 mm', confirming the fix with a second oracle.

Silkscreen legibility (owner requirement) - 95 mis-attributed labels -> 3

metric	before	after
median refdes offset	4.079 mm	1.600 mm
>1 mm beyond own pad extent (excl. H1-H5)	95	3
>1.5 mm beyond own pad extent	86	1
median beyond-extent	2.769 mm	0.119 mm
visible refdes	116/116	116/116

Root cause was a **library** defect, not the P6 sweep alone: 'easyeda2kicad' emits every

footprint's Reference at a blanket '(0, -4.0)' mm regardless of part size, and the large

packages instead derive theirs from the silk outline (WROOM-1 at -12.955 mm).

The library was

normalized centrally (31/31 footprints, idempotent, 'lib/EDITS.md' edit 8); the board needed

111 'move_text' ops because it carries its own copies. Applied in 5 batches, DRC-verified

after each. Non-regression proven by snapshot diff: 3294 tracks, 359 vias, 8 zones, every pad,

every footprint silk and every 'value' field byte-identical; 0 footprints moved.

****Attribution beat closeness on 17 refs****; 110 of 111 have their own part as nearest.
 One residual: ****C35 at 3.861 mm****, boxed in on four sides, whose only wide-enough channel is
 C34's only attributable slot - a P6 placement consequence, not a silk problem.

The measurement lesson of this phase

****‘check_creepage’ cannot express a coated board.**** It hardcodes IPC-2221 rows B2 (external ****uncoated****) and B1 only. Re-adjudicated per item type against the correct rows - B4 for a mask-covered conductor (****0.13 mm**** at 51-100 V), A6 for an exposed land (****0.50 mm****, and 0.80 mm at 101-150 V per IPC 6.3.4), B1 for inner layers - the count goes ****11 -> 10****, and the one that clears is the largest: `‘/poe/POE.TAP_A2 <-> /poe/LED.Y.A’`, whose ****217 pairs**** under 0.600 mm are all trace-to-trace under soldermask and pass B4 with +0.0731 mm. Two bounded fixer attempts and one applied-then-retracted edit were spent on a requirement that does not apply to that geometry. The remaining 10 are every one of them a pad inside a single package (U1 x7, U22 x2, U20 x1) and fail every applicable row, because 0.200 mm is the inter-lead gap of the SOIC-8 that TI ships.

Also worth carrying forward

`‘gate dfm’` already reads ****194 errors**, 189 of them "trace width 0.1000 mm below JLC minimum 0.1016 mm"******. Root cause is a configuration defect from P5: `‘board.init’` writes `‘min_track.width: 0.1’` - below ****every**** JLC profile - and this board's hand-written `‘.kicad_dru’` contains none of `‘rules_gen’`'s generated floors. So `‘drc_routed’` 0/0 currently does ****not**** imply fabricable, and the gap is 1.6 micrometres. This is a P9 item, reported not fixed.

Gate attempt history

Timestamp	Gate	Status	Failing/Total
2026-07-28T20:16:22	erc	pass	0/0
2026-07-28T20:56:57	erc	pass	0/0
2026-07-29T11:02:27	place	pass	0/0
2026-07-29T17:21:37	drc_routed	pass	0/0
2026-07-29T17:25:18	verify	fail	220/224
2026-07-29T18:55:29	drc_routed	pass	0/0
2026-07-29T22:11:27	verify	fail	114/118

Issues

(no entries)

Human checkpoints

Id	Phase	Status	When	Note
1	P2	approved	2026-07-28T15:00:34	Q5 plastic enclosure + heatsink enclosed, non-isolated stands (closes open item C). Q8 Wi-Fi FUNCTIONAL - overrides carrier brief s6; antenna keepout mandatory on every layer, RF review in scope at layout sign-off. Q4 11.0 mm stack + 30x26 mm daughter notch (anti-180 interlock). Q6 adopt derived numbers: 48V 0.25A af / 0.5A at sustained, 1.80 A/pin, eFuse 1.0 A, J3 stays 2x7. Q7 module WROOM-1-N8.
2	P4	approved	2026-07-28T21:27:15	Magjack: qualify a Bel/Pulse PoE+ ICM off-LCSC (owner rejected af-only for consistency with the par emitter rejection). Q9: Ethernet OTA + Wi-Fi OTA + recovery header, no USB-C, no pogo jig. OVP: 0.1 percent resistors on R66/R67/R73.
3	P6	not reached		
4	P8	not reached		
5	P10	not reached		

10 Artifact Index

File	Size	Note
state.json	95,849 B	
kicad/lumina-carrier.kicad_pcb	2,078,378 B	
kicad/lumina-carrier.kicad_sch	40,312 B	
architecture/blocks.md	32,381 B	
architecture/connector-icd.md	50,471 B	
architecture/constraints.json	8,733 B	
architecture/decisions.md	22,677 B	
architecture/power_tree.md	25,946 B	
architecture/requirements.md	27,840 B	
architecture/sheets.md	17,437 B	
architecture/stackup.md	21,413 B	
reports/drc-place.json	258,020 B	
reports/erc-waivers.md	9,722 B	
reports/fp_D10_C19829453.overlay.svg	2,597 B	
reports/fp_J1_C91754.overlay.svg	7,027 B	
reports/fp_U10_C32843.overlay.svg	18,050 B	
reports/fp_U1_C337500.overlay.svg	3,690 B	
reports/fp_U1_C337500.kicadstock.overlay.svg	3,733 B	
reports/fp_U20_C5124114.overlay.svg	3,692 B	
reports/fp_U21_C116592.overlay.svg	2,597 B	
reports/fp_U22_C1849461.overlay.svg	9,647 B	
reports/fp_U30_C2913198.overlay.svg	22,797 B	
reports/fp_verify_D10_C19829453.json	1,570 B	
reports/fp_verify_J1_C91754.json	2,167 B	
reports/fp_verify_U10_C32843.json	2,101 B	
reports/fp_verify_U1_C337500.json	2,872 B	
reports/fp_verify_U1_C337500.kicadstock.json	3,624 B	
reports/fp_verify_U20_C5124114.json	1,682 B	
reports/fp_verify_U21_C116592.json	1,482 B	

File	Size	Note
reports/fp_verify_U22_C1849461.json	2,840 B	
reports/fp_verify_U30_C2913198.json	3,148 B	
reports/gate-place.json	337 B	
reports/lib_pull.json	14,733 B	
reports/lumina-carrier-schematic.pdf	483,671 B	
reports/place_anneal.json	5,013 B	
reports/place_metrics_final.json	11,906 B	
reports/place_region_audit.json	3,500 B	
reports/place_seed.json	7,846 B	
reports/route-probe-p6.json	326 B	